



#MDBlocal



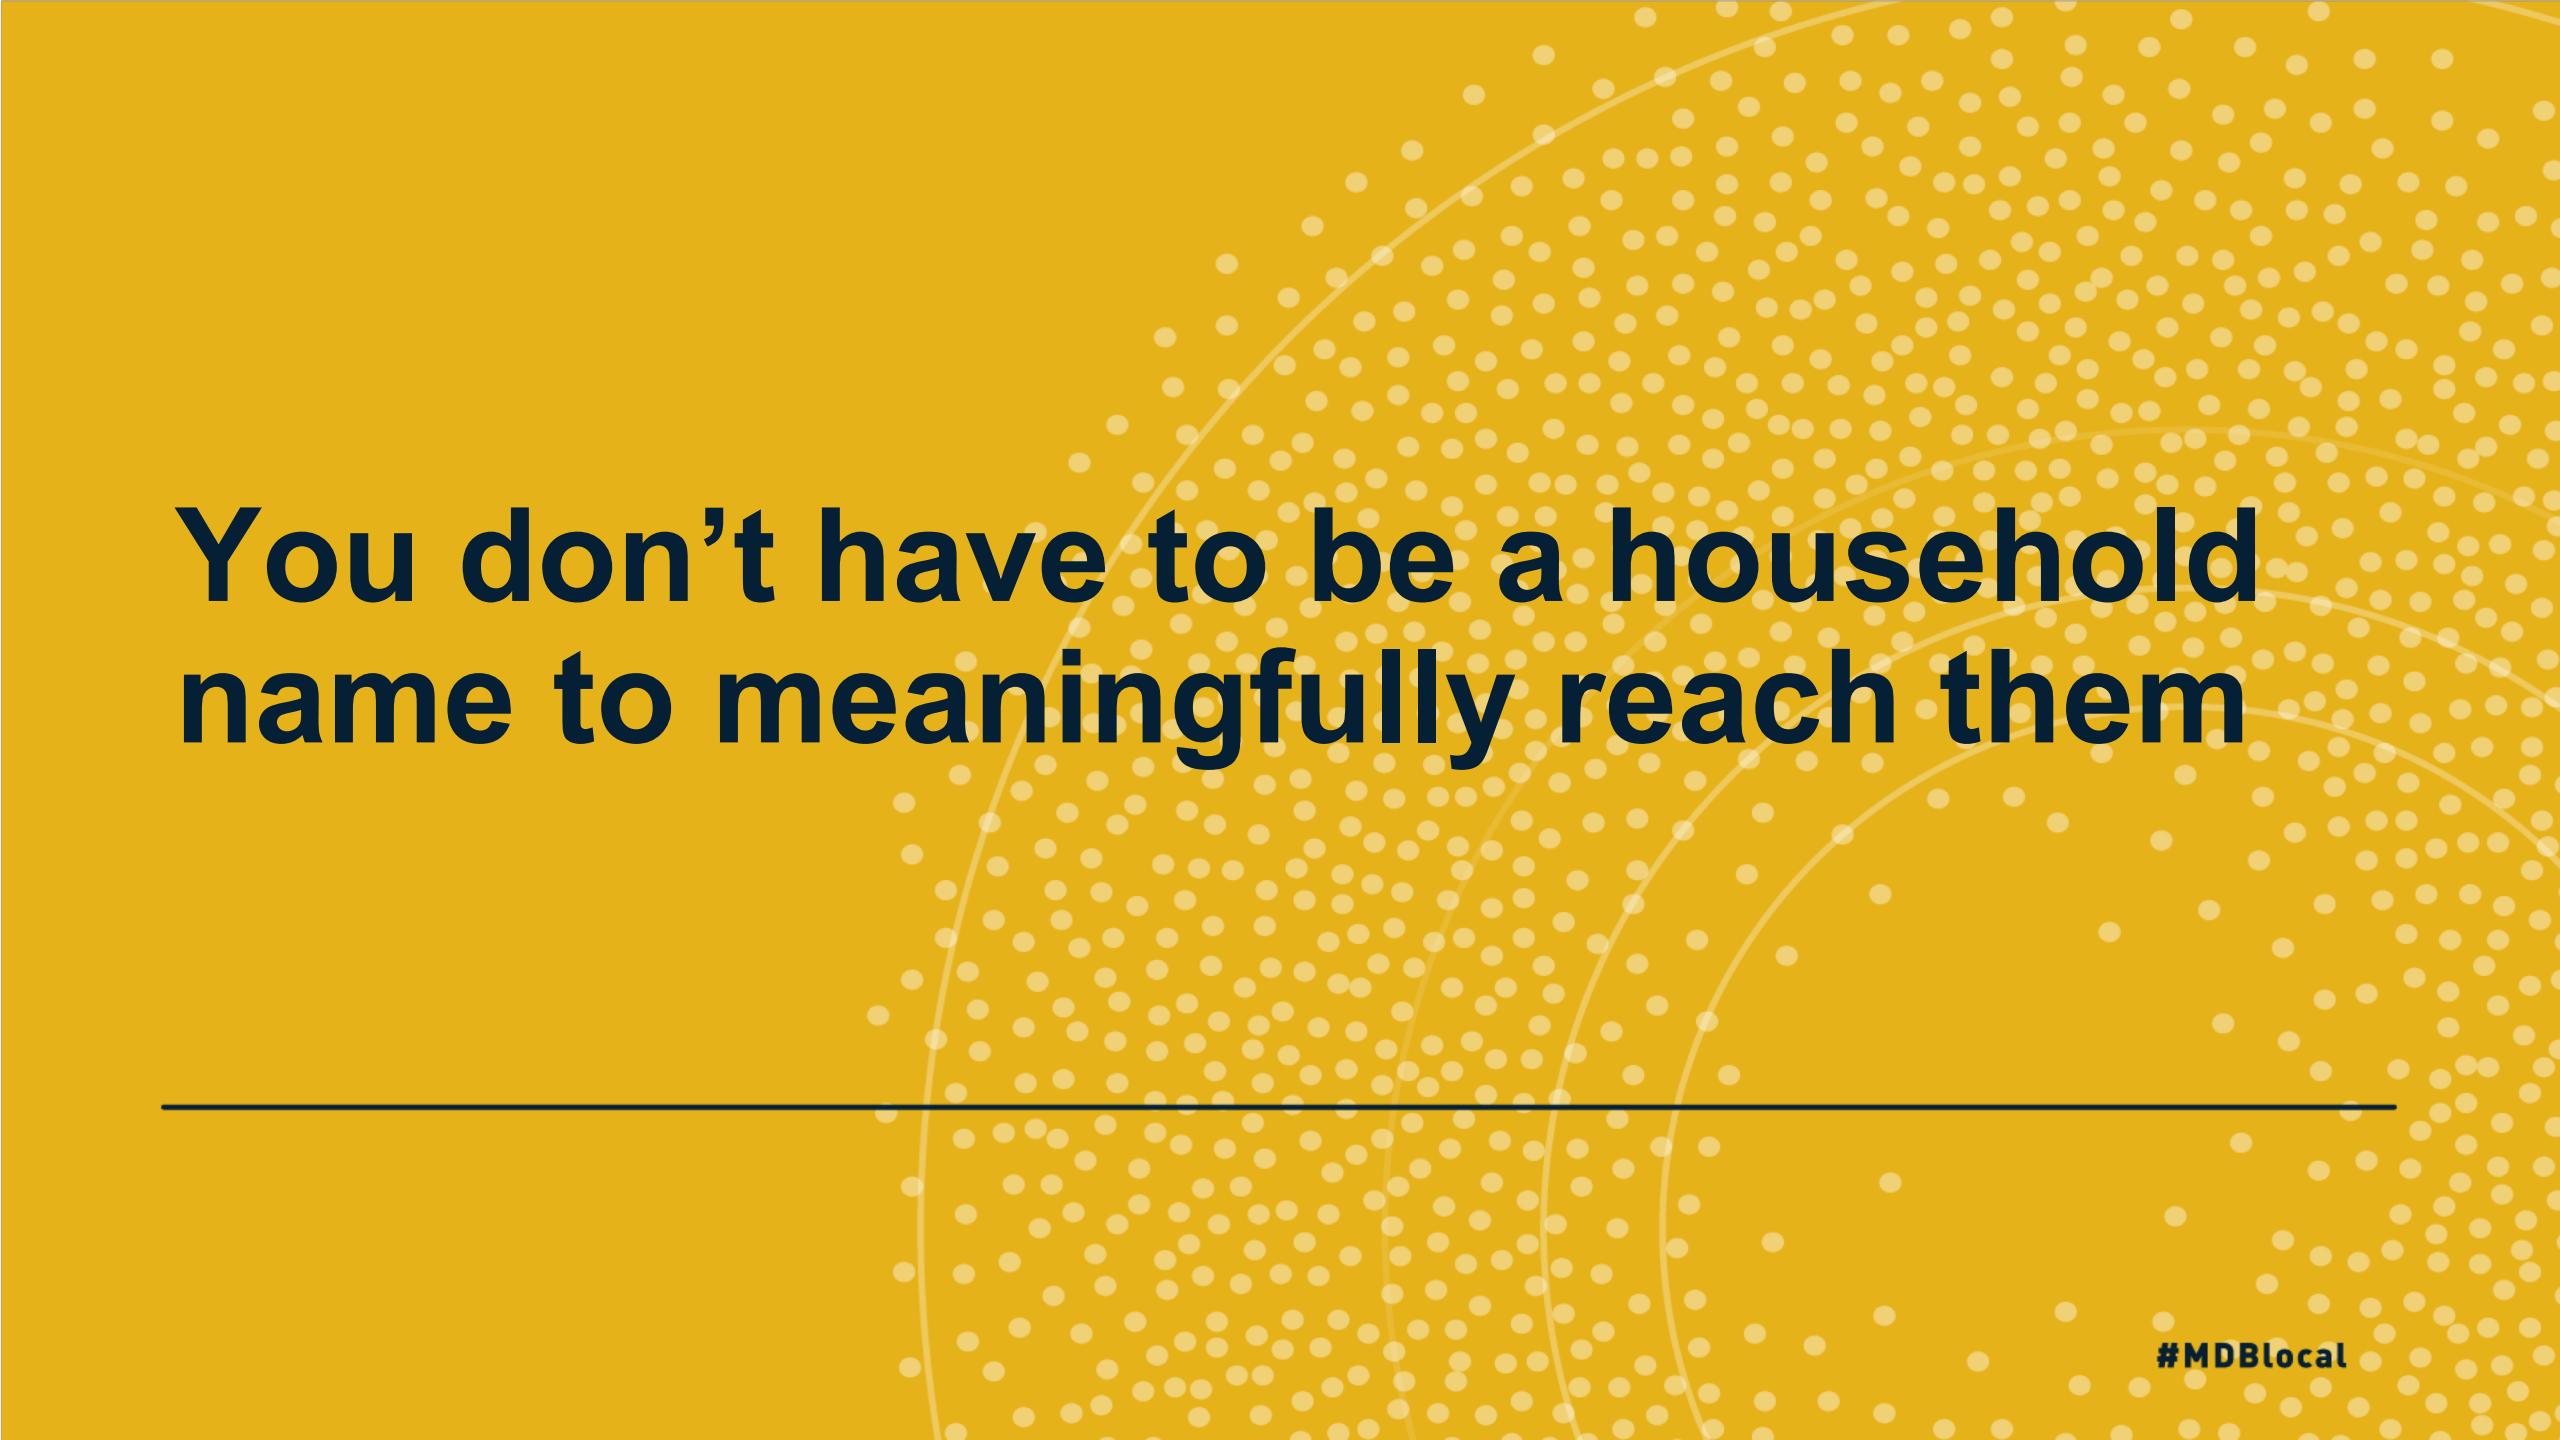
Global Clusters Topologies

Sig Narvaez

Principal Solution Architect, MongoDB

3.15 Billion
new internet users
since the first MongoDB code commit

**Every year a US-sized population
is coming online for the first time**



You don't have to be a household name to meaningfully reach them

**Yes, a CDN can *reach*
but users expect *more***

Global Operational Data

Your users want...

Interactive experiences

Personalization

Instant gratification

Global Operational Data

Your users want...

Interactive experiences

Personalization

Instant gratification

You want...

Happy users

To stay ahead of competition

Global awareness

What does **100ms** equate to?

What does **100ms** equate to?

1% drop in sales

**Missed chance to engage a customer
meaningfully**

Global Latencies from AWS US West

Americas	US West – US East:	83ms
EMEA	US West – Frankfurt:	163ms
LATAM	US West – Sao Paulo:	190ms
APAC	US West – Singapore:	225ms

Options

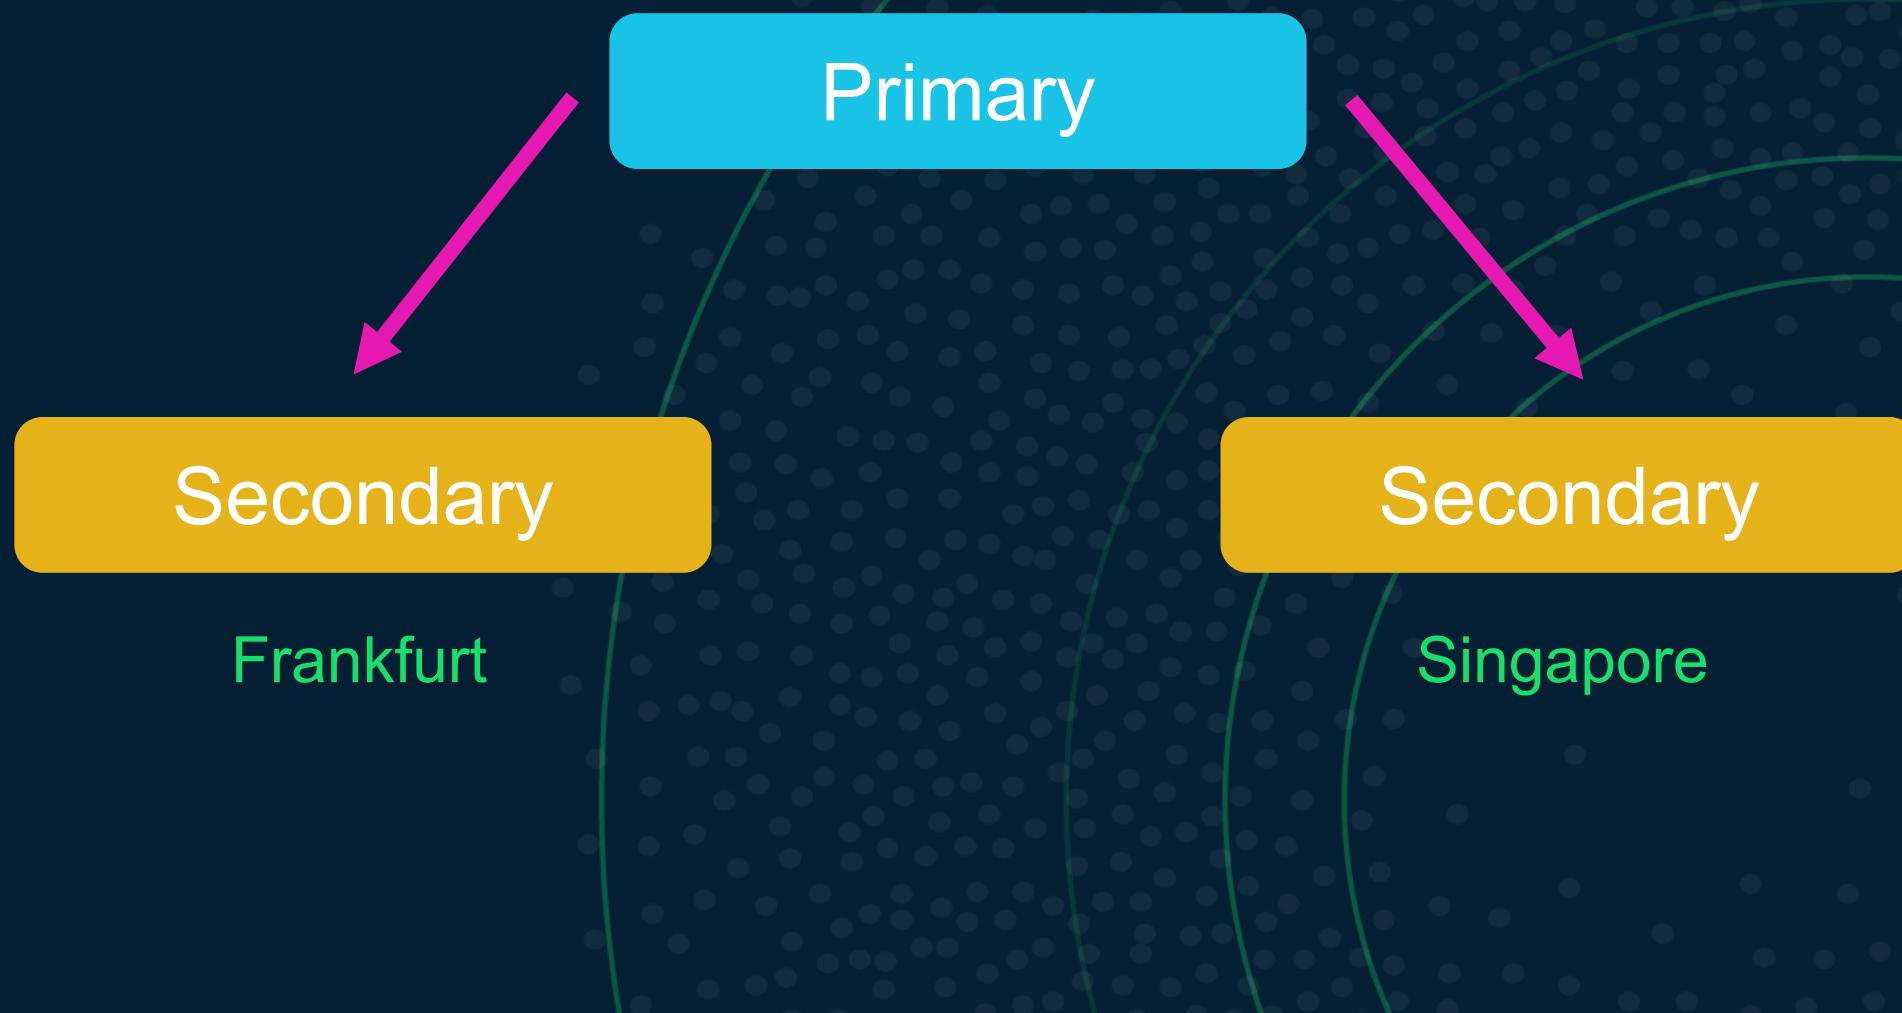
Global Multi-Region Replication

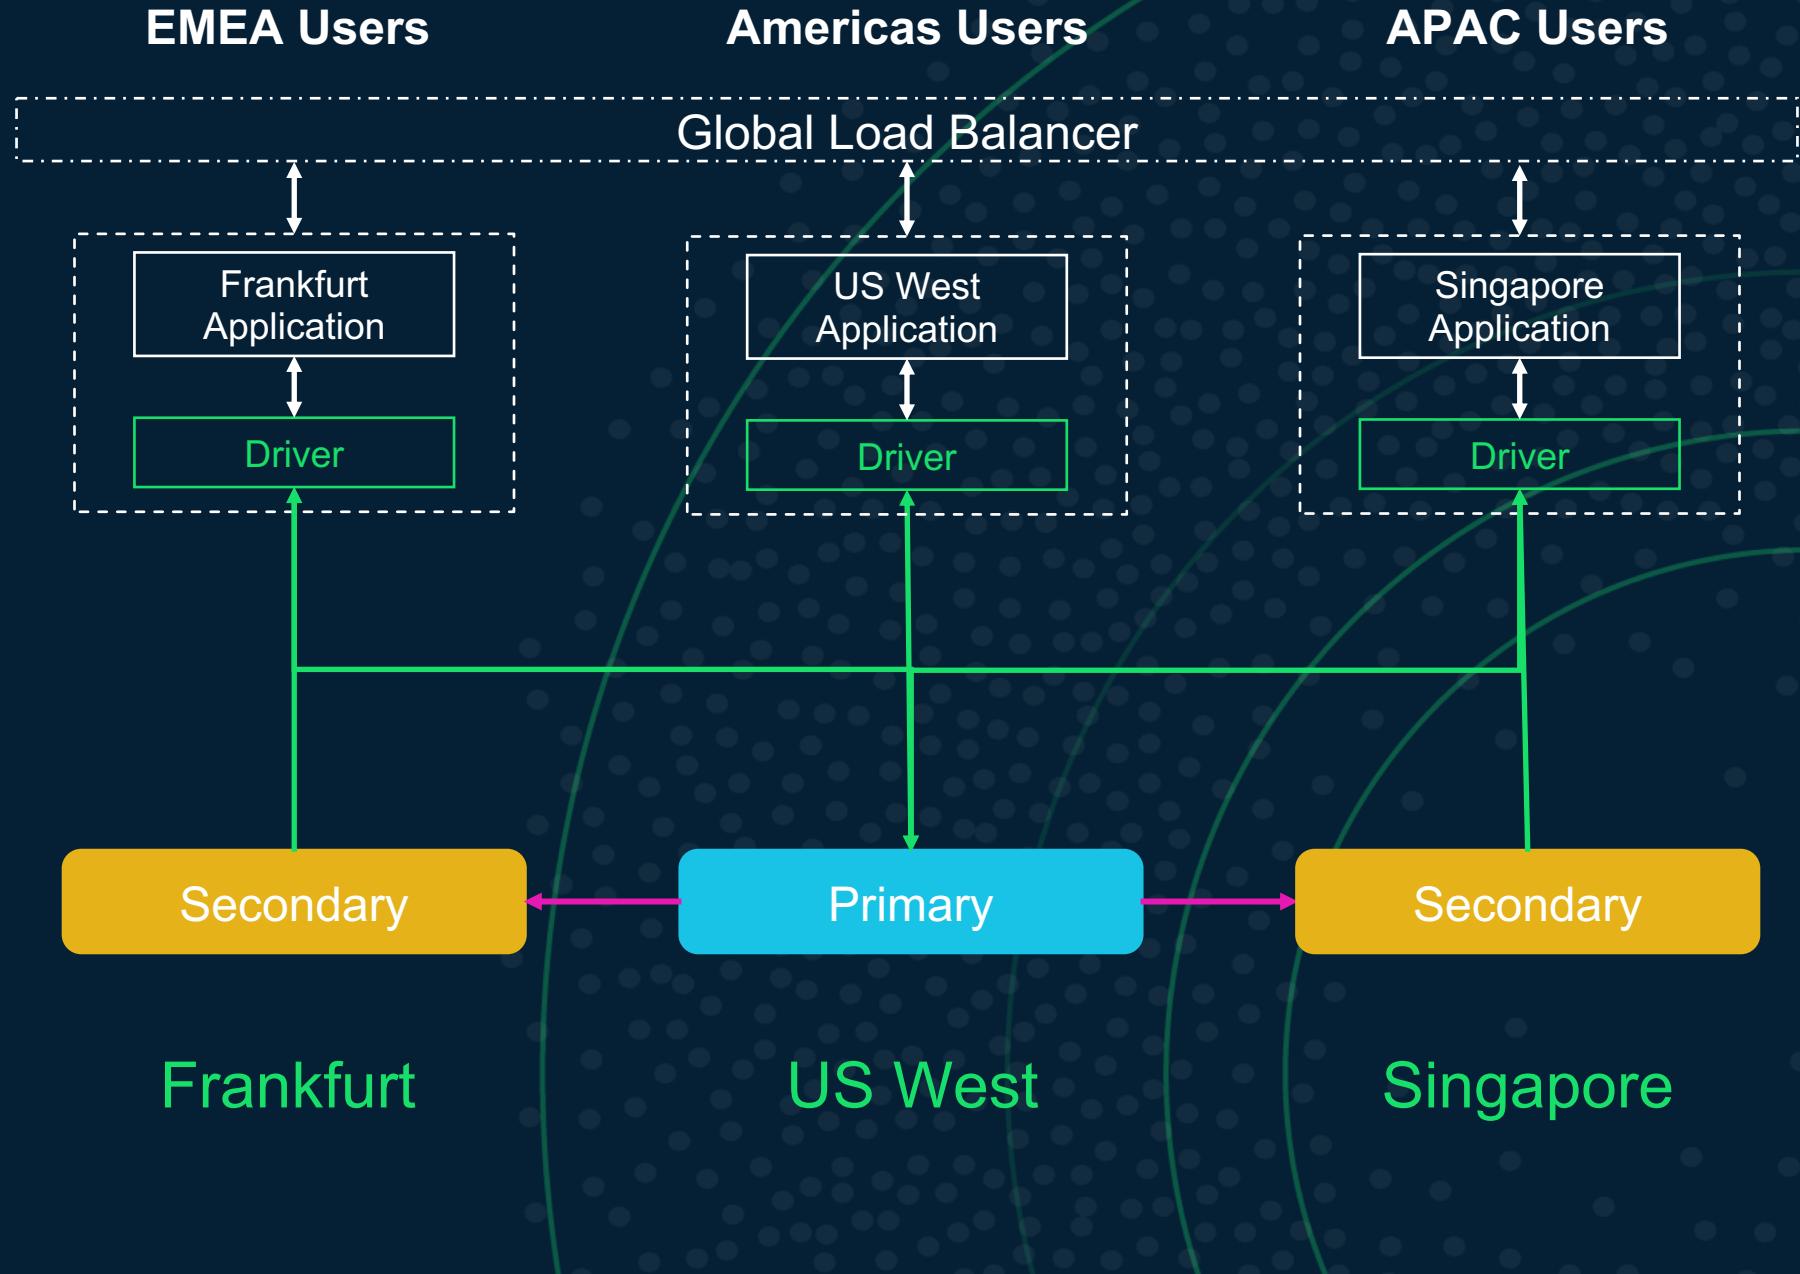
Primary

Secondary

Secondary







Global Multi-Region Replication

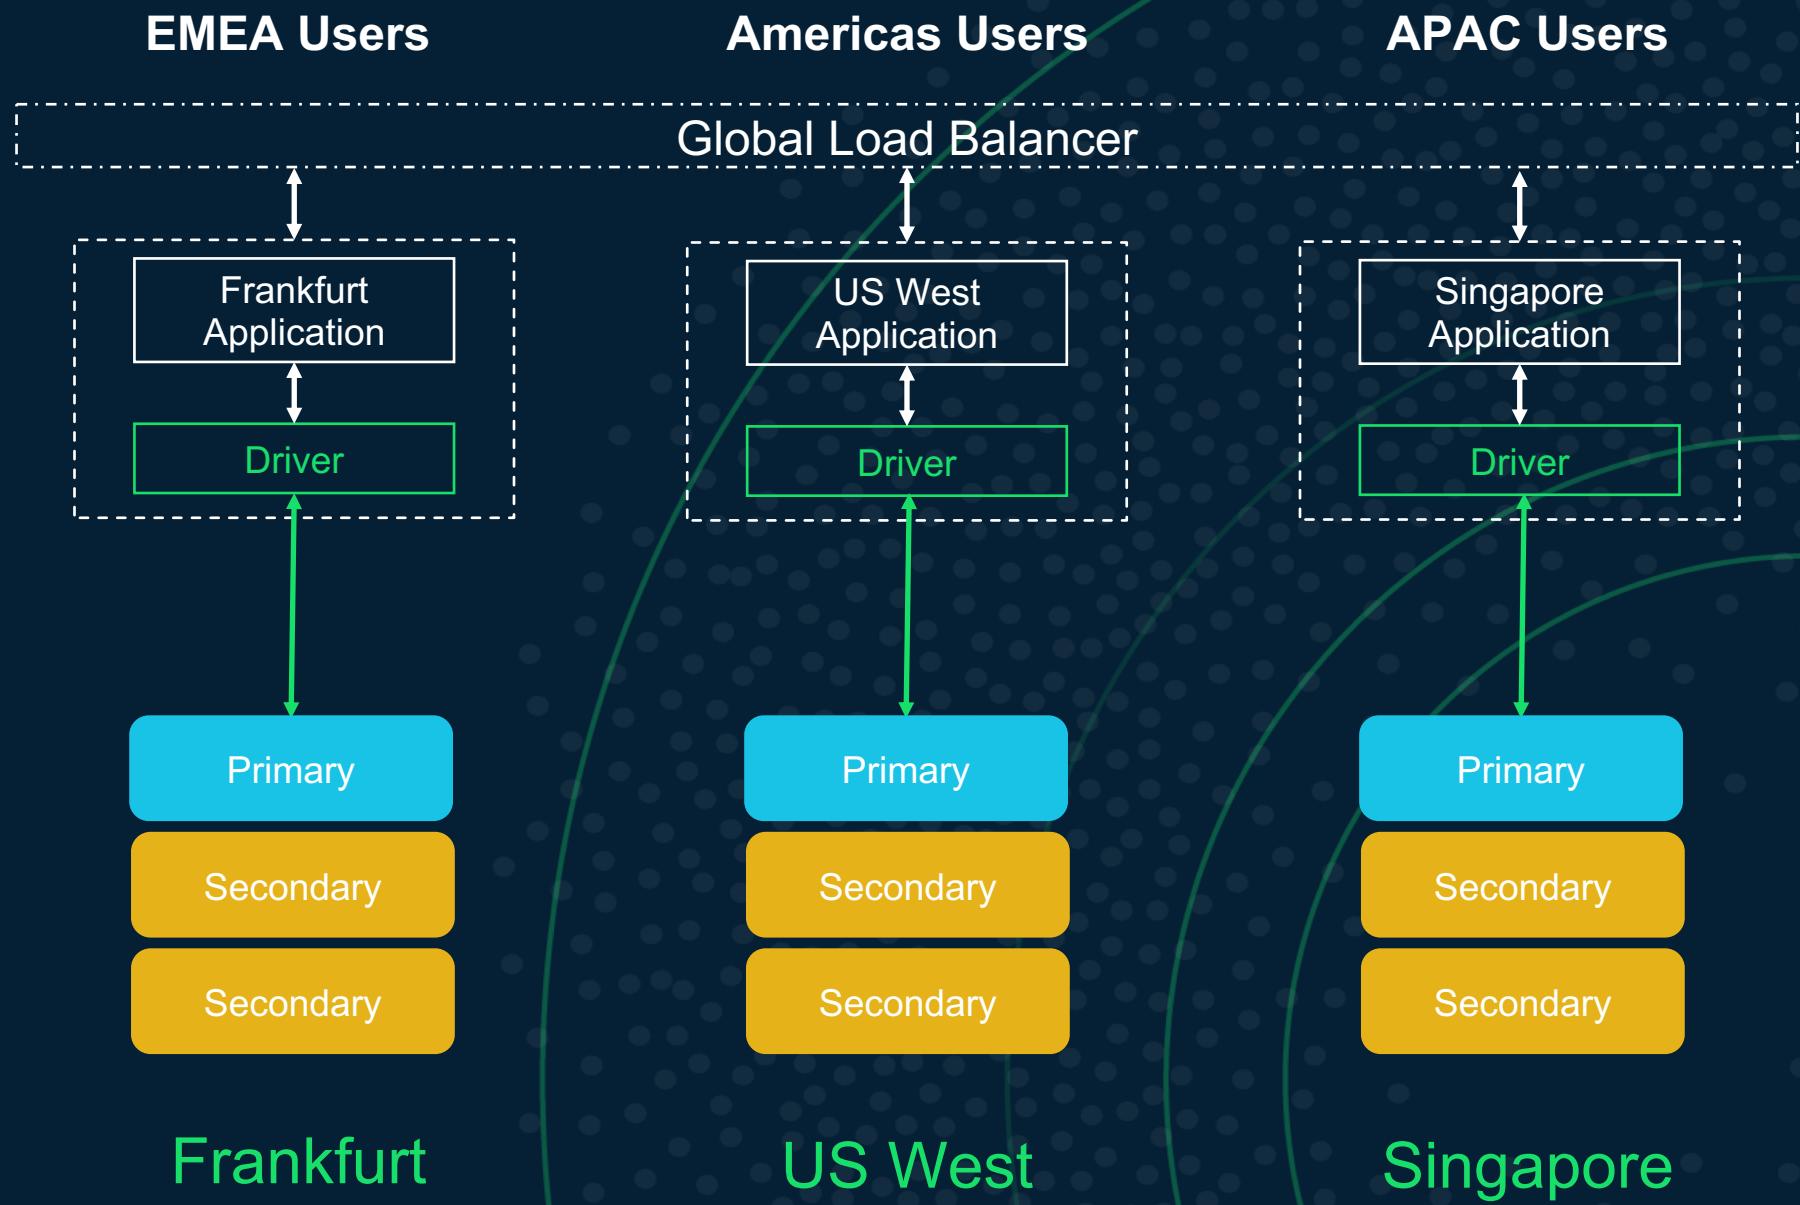
Advantages

Single connection string globally
Cross-region high availability

Disadvantages

All writes go to preferred region
 $\frac{2}{3}$ of users experience $> 100\text{ms}$
Fast reads require secondary read preference
Asymmetrical deployment

Cluster Per Region



Cluster Per Region

Advantages

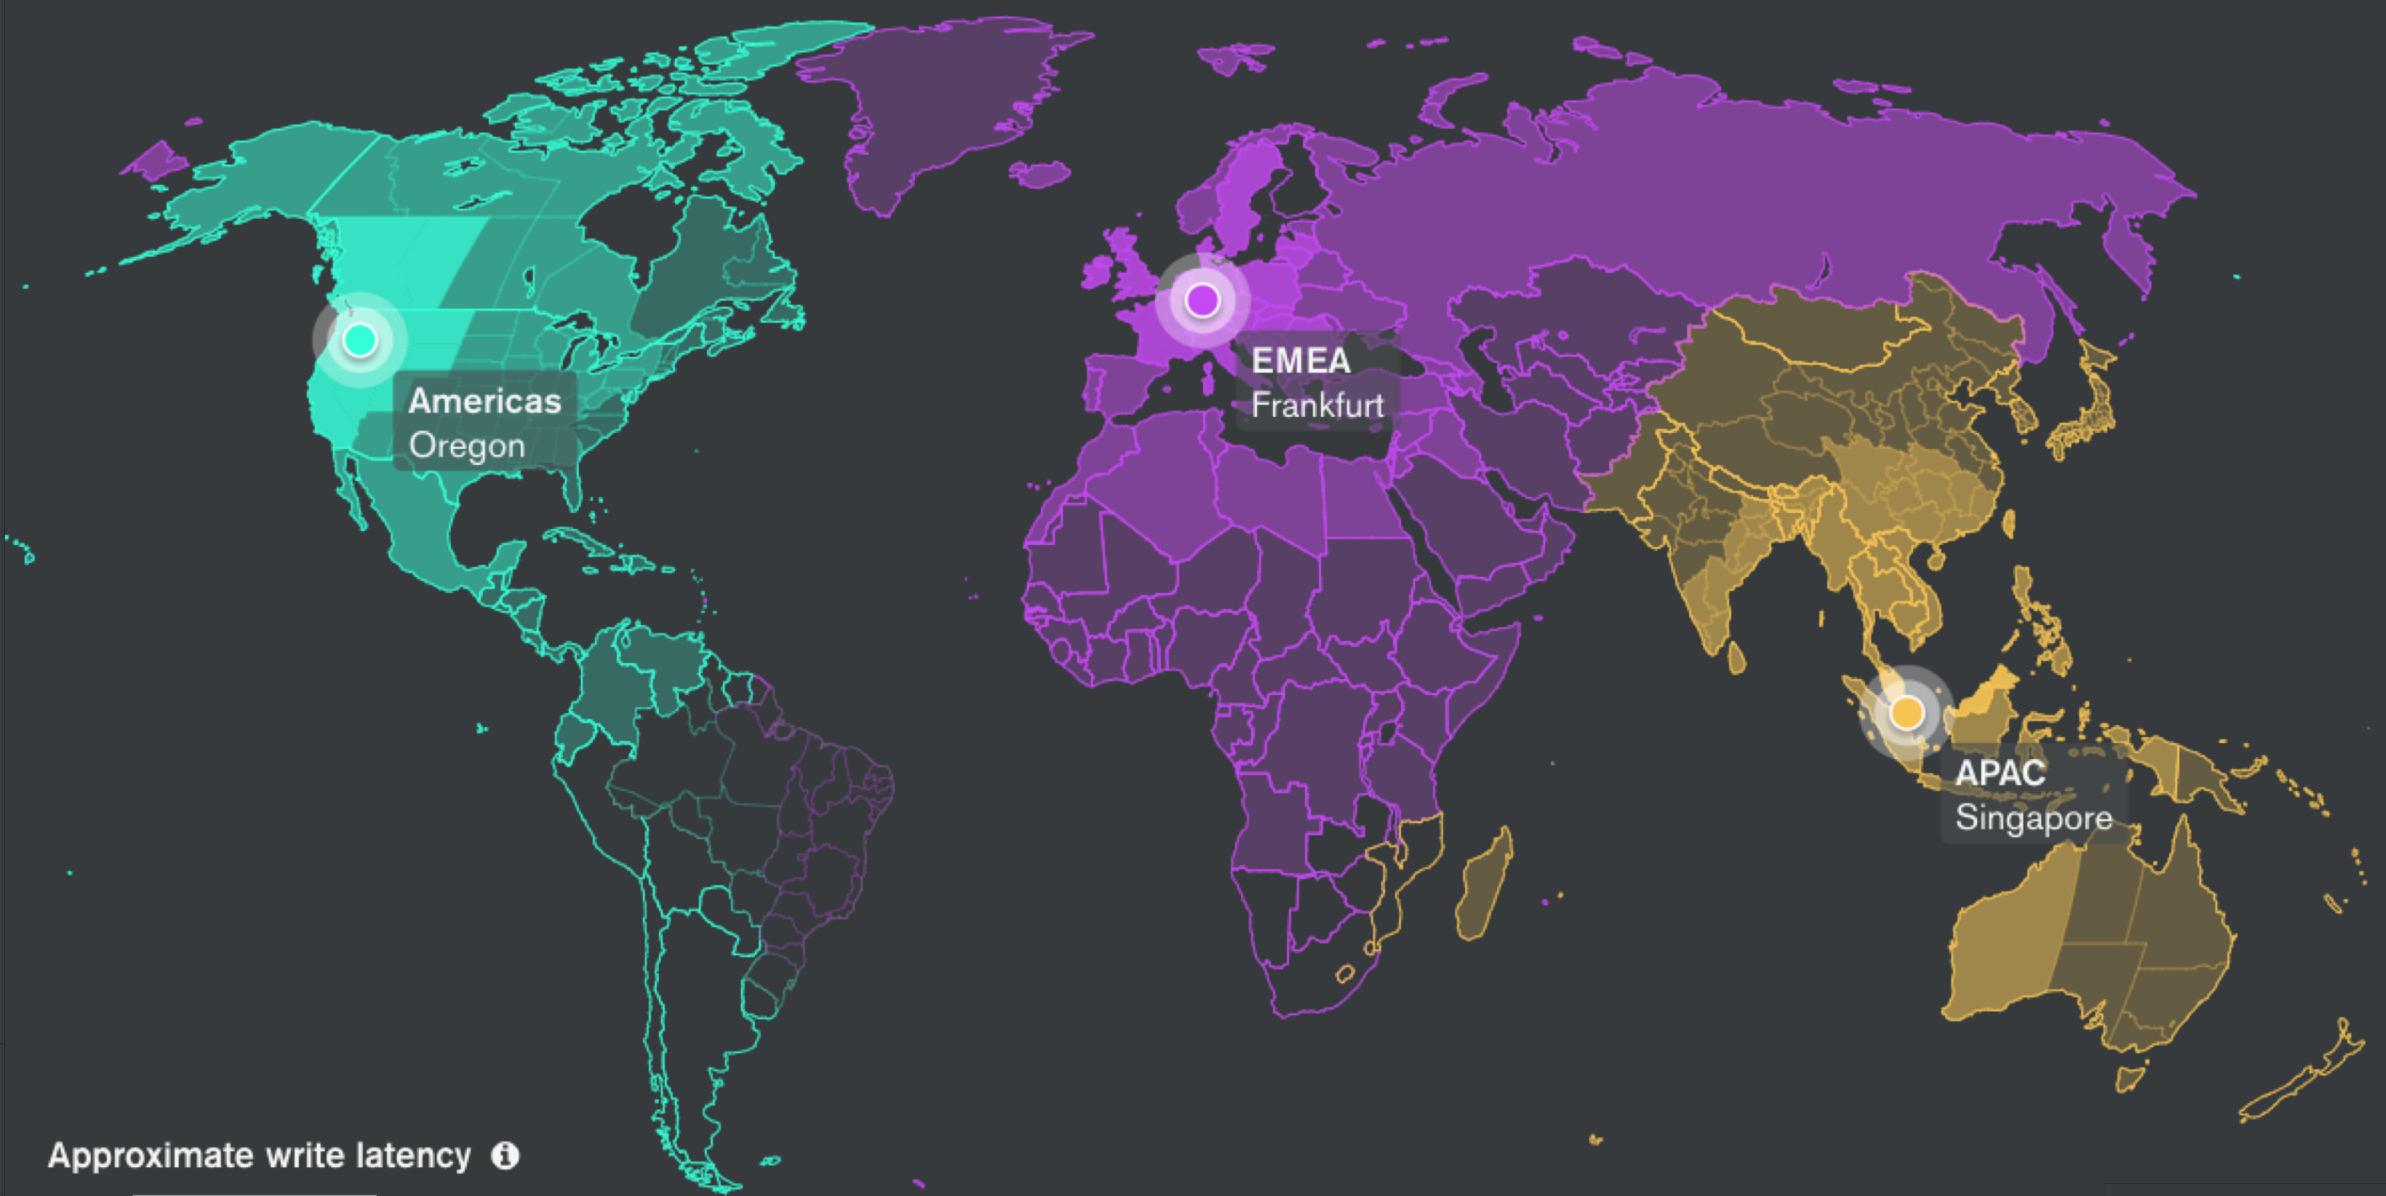
Easy today in MongoDB Atlas
Independently scalable

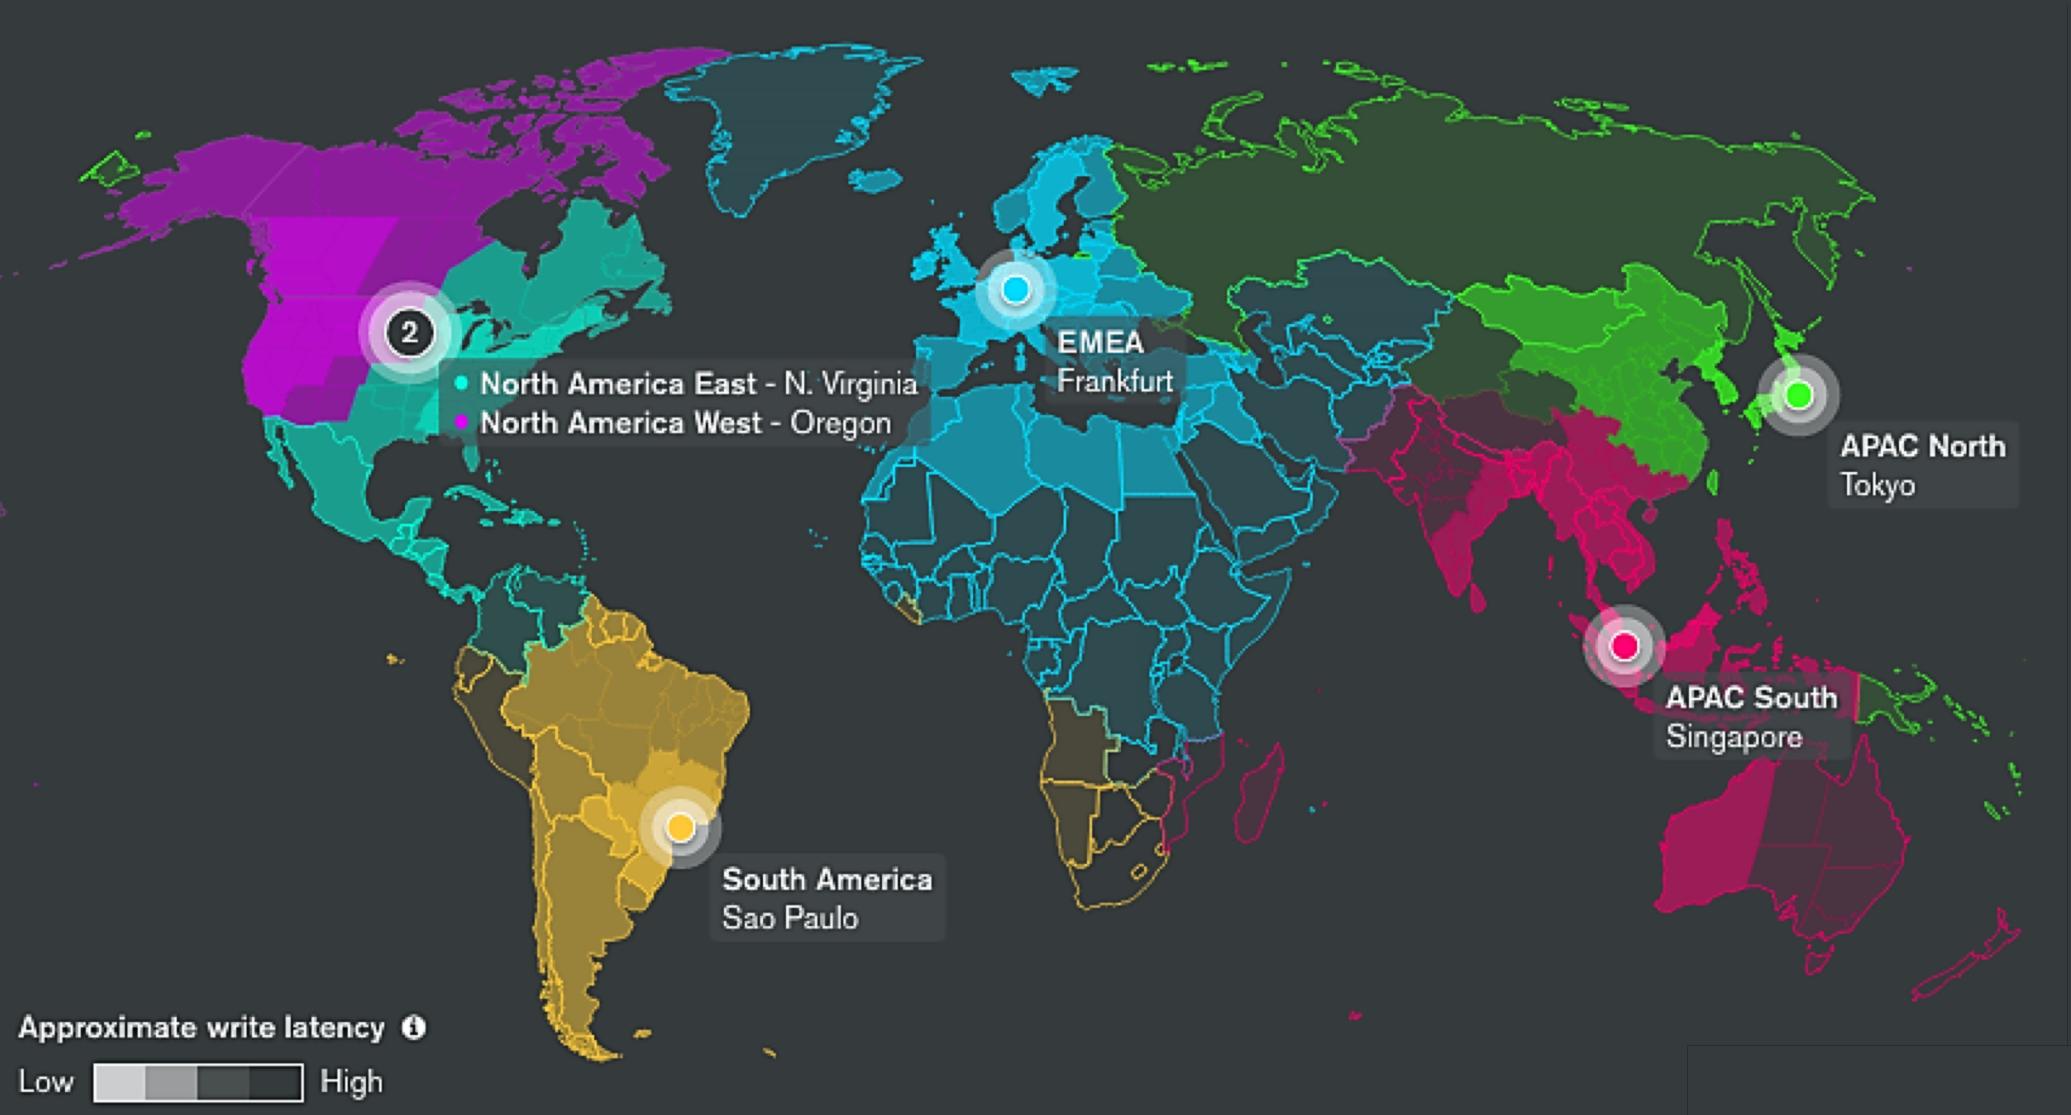
Disadvantages

No global aggregation
Application tier manages end user to connection string affinity

Can we get the best of both worlds?

Global Clusters with Global Writes





How Global Writes work

Data includes location attribute

```
{  
  _id: 123,  
  firstName: "Johannes",  
  lastName: "Doe",  
  activities: [  
    {  
      title: "Skiing",  
      attributes: ["downhill", "jumps"]  
    },  
    {  
      title: "Coding",  
      attributes: ["mongodb", "python"]  
    }  
  ]  
}
```



```
{  
  _id: 123,  
  location: "DE",  
  firstName: "Johannes",  
  lastName: "Doe",  
  activities: [  
    {  
      title: "Skiing",  
      attributes: ["downhill", "jumps"]  
    },  
    {  
      title: "Coding",  
      attributes: ["mongodb", "python"]  
    }  
  ]  
}
```

What is a location attribute?

Generally...

an ISO 3166-1 country code

e.g., US, DE, IN, or JP

Or optionally and more specifically...

an ISO-3166-2 subdivision code

e.g., US-CA or US-NY

Examples	Description	Potential Region Affinity
US	United States	US East
DE	Germany	Frankfurt
IN	India	Mumbai
JP	Japan	Tokyo
US-CA	California	US West
US-NY	New York	US East

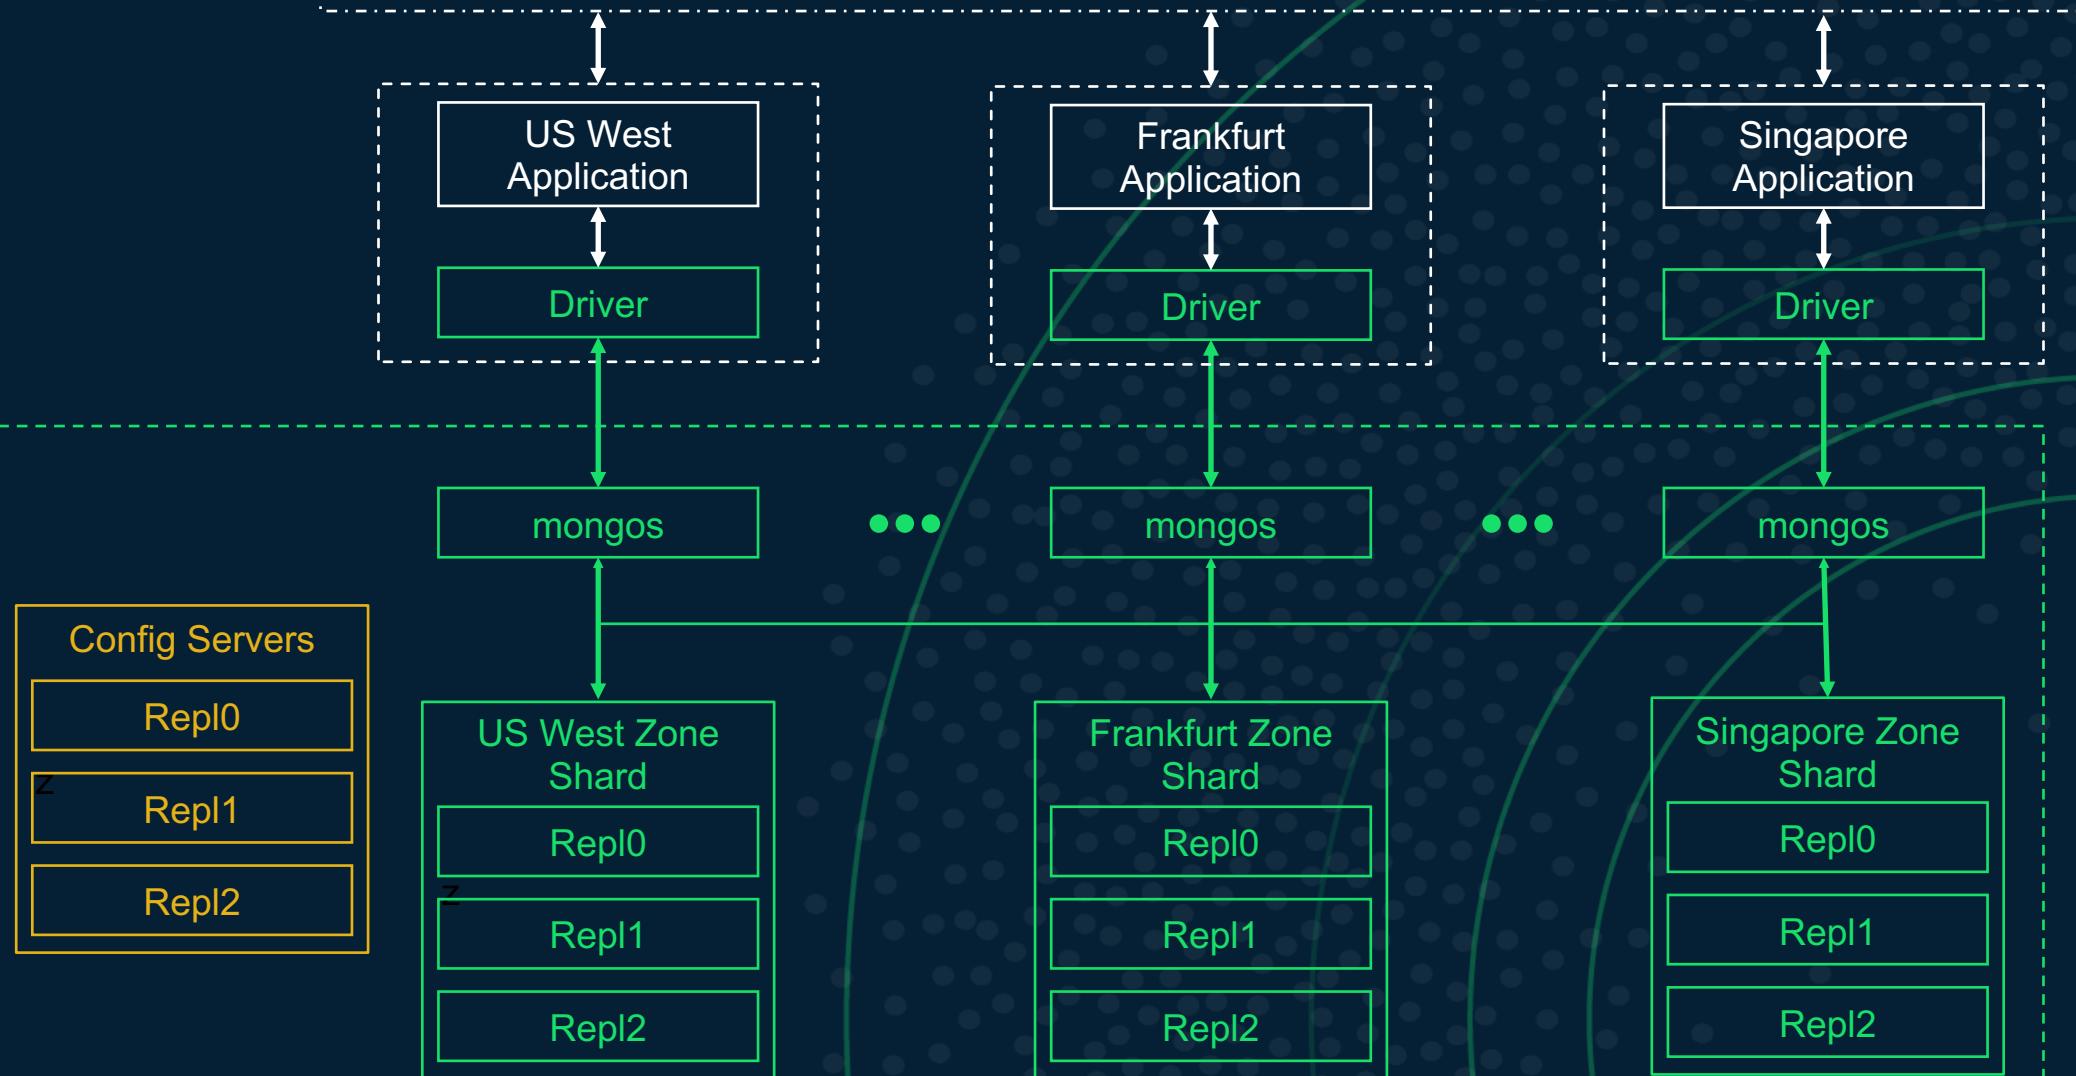
Americas Users

EMEA Users

APAC Users

Global Load Balancer

MongoDB Atlas Global Cluster



Data includes location attribute

```
{  
  _id: 123,  
  firstName: "Johannes",  
  lastName: "Doe",  
  activities: [  
    {  
      title: "Skiing",  
      attributes: ["downhill", "jumps"]  
    },  
    {  
      title: "Coding",  
      attributes: ["mongodb", "python"]  
    }  
  ]  
}
```



```
{  
  _id: 123,  
  location: "DE",  
  firstName: "Johannes",  
  lastName: "Doe",  
  activities: [  
    {  
      title: "Skiing",  
      attributes: ["downhill", "jumps"]  
    },  
    {  
      title: "Coding",  
      attributes: ["mongodb", "python"]  
    }  
  ]  
}
```

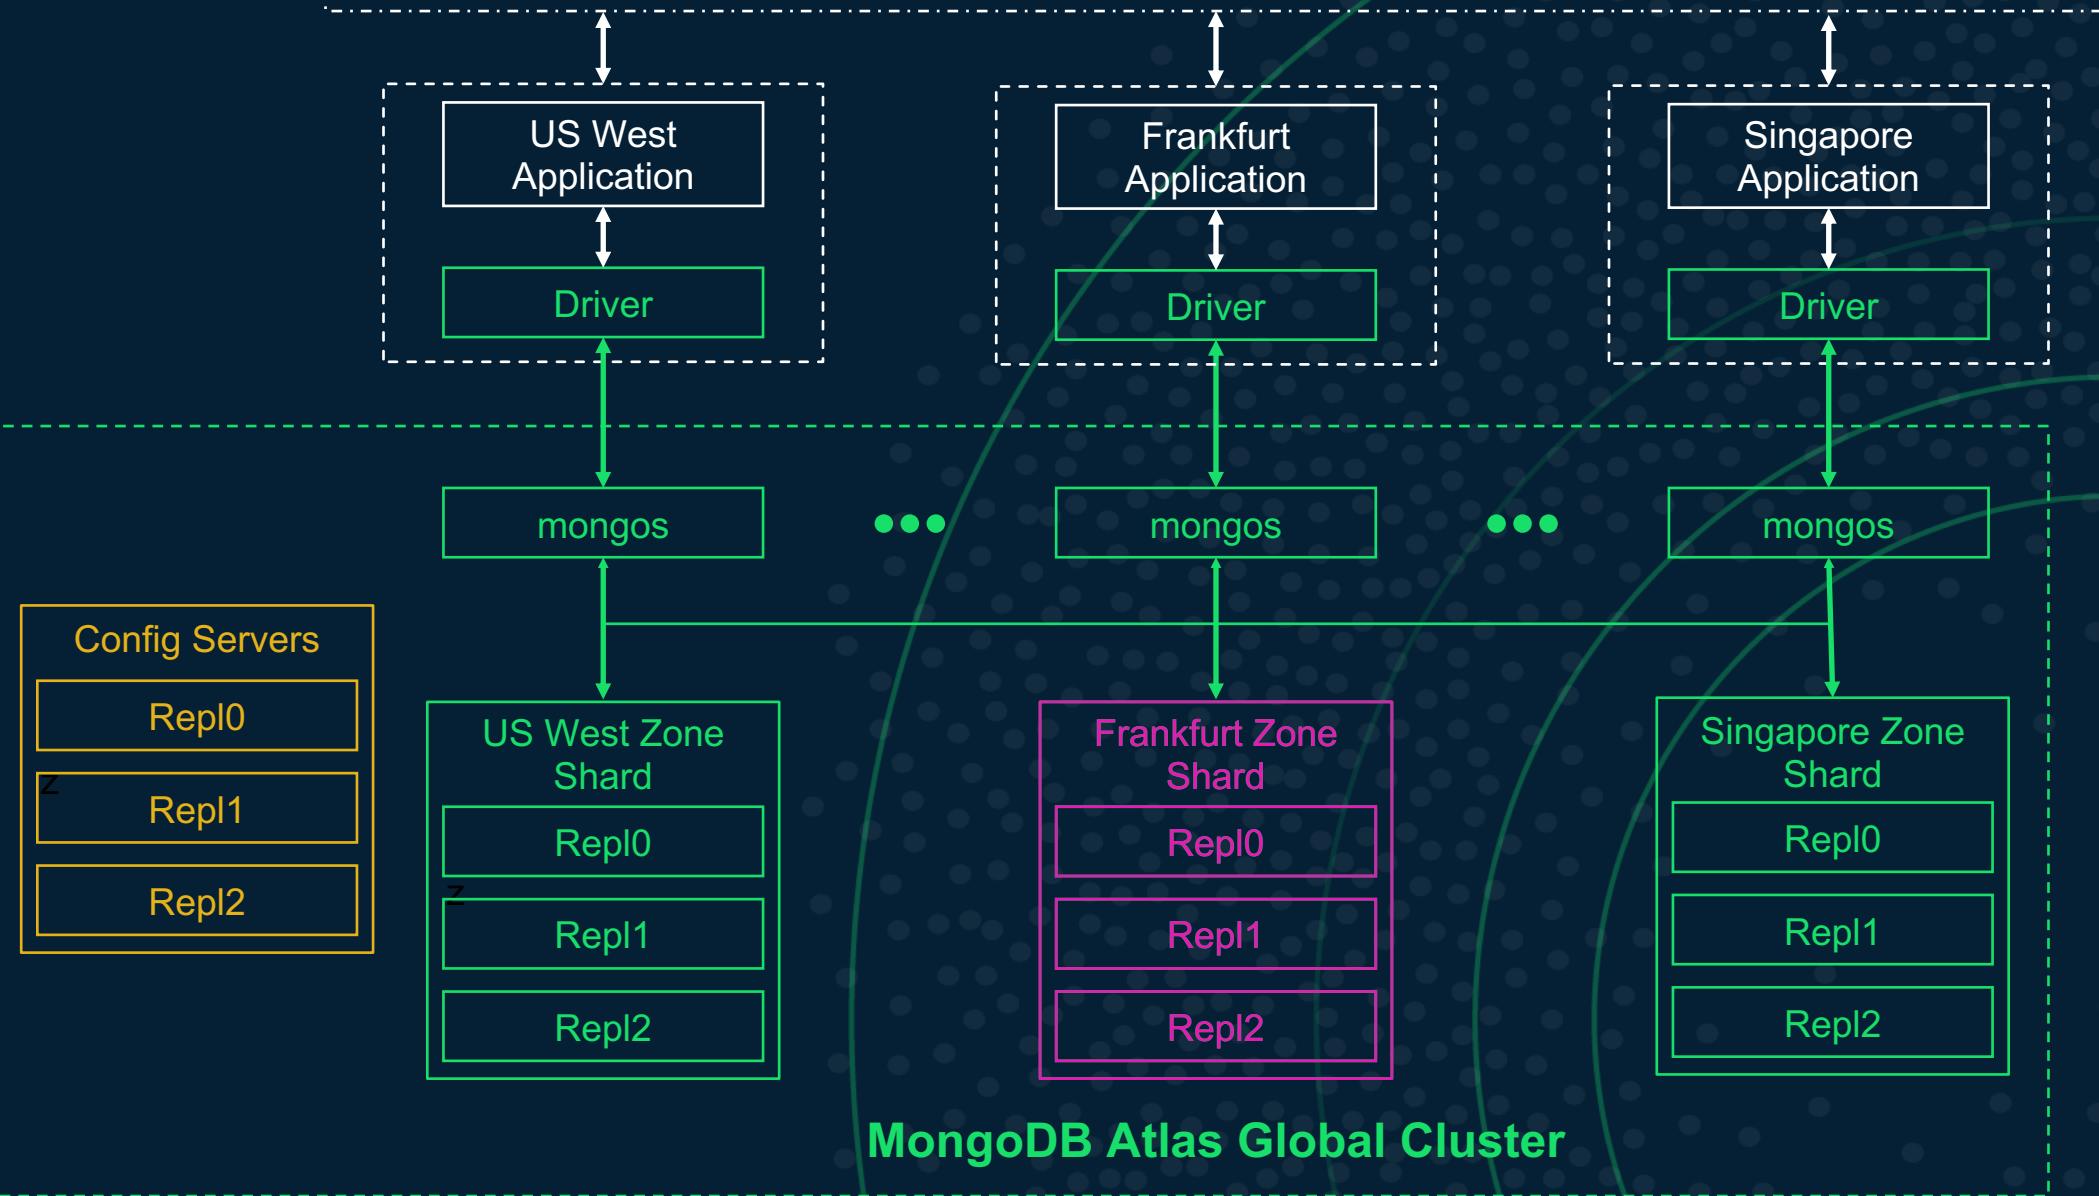
Americas Users

EMEA Users

APAC Users

Global Load Balancer

MongoDB Atlas Global Cluster



Define Global Collections

MongoDB Atlas will be very opinionated...

In the Data Explorer, you'll be guided to set a shard key like

```
{location : 1, <identifier> : 1}
```

You'll receive global-targeting query examples

```
db.globalColl.findOne({ location : "XY", <identifier> : 123, ... })
```

For example location and `_id`

```
{  
  _id: 123,  
  location: "CA",  
  firstName: "Jane",  
  lastName: "Doe",  
  activities: [  
    {  
      title: "Mountaineering",  
      attributes: ["glaciers"]  
    },  
    {  
      title: "Coding",  
      attributes: ["java", "mongodb"]  
    }  
  ]  
}
```

The screenshot shows the MongoDB Compass interface. At the top, it says "Cluster0" with tabs for "Overview", "Metrics", "Data" (which is selected), and "Command Line Tools". Below that, it shows "2 DATABASES 2 COLLECTIONS" and a button "+ Create Database". There's a search bar "NAMESPACES" and a tree view of databases and collections: "globalTest" (expanded) containing "globalColl" and "testdb", and "testdb" (collapsed). On the right, under "globalTest.globalColl", it says "COLLECTION SIZE 18.9MB" and has tabs for "Find" (selected) and "Indexes". A "FILTER" button is shown with the query: {"filter": {"location": "CA"}}, and below it, "QUERY RESULTS 1-20 OF 1" is displayed. The result is a single document: {_id: ObjectId("596f2d1a1a1a1a1a1a1a1a1a"), date: "2017-06-01", location: "CA"}.

How are locations mapped to zones?

Configure Zone Mapping

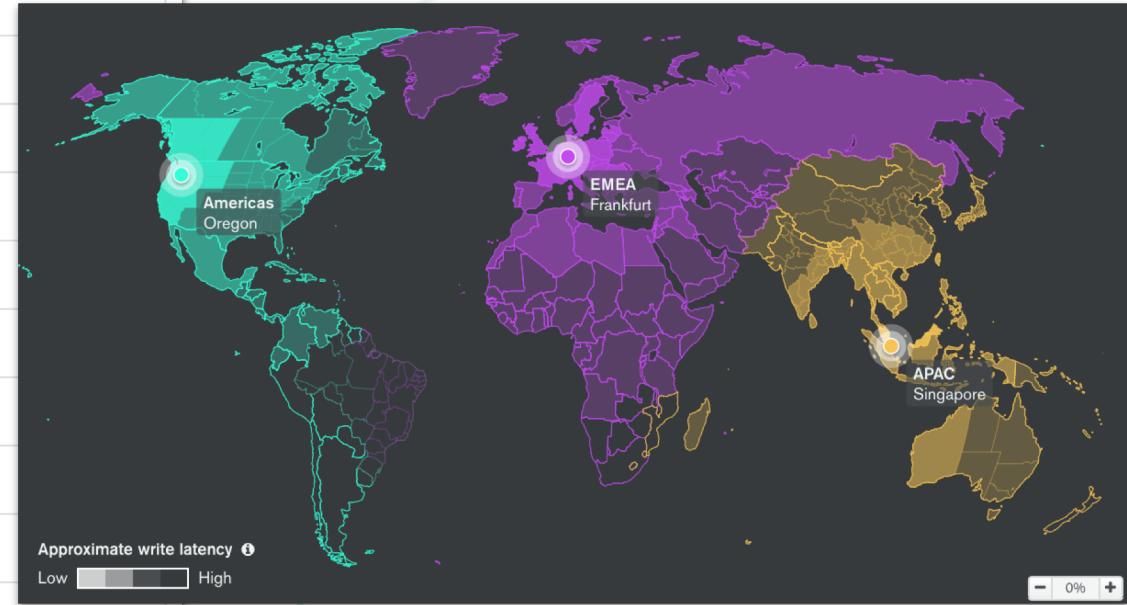
We automatically map every location in the world to the zone with the closest preferred region.

Search for a location 

[BACK TO MAP](#)

Bhutan (BT)	APAC (Singapore)
Bolivia, Plurinational State of (BO)	Americas (Oregon)
Bonaire, Sint Eustatius and Saba (BQ)	Americas (Oregon)
Bosnia and Herzegovina (BA)	EMEA (Frankfurt)
Botswana (BW)	EMEA (Frankfurt)
Bouvet Island (BV)	APAC (Singapore)
▼ Brazil (BR) 27 SUBDIVISIONS	EMEA (Frankfurt)
Brazil - Acre (BR-AC)	Americas (Oregon)
Brazil - Alagoas (BR-AL)	EMEA (Frankfurt)
Brazil - Amazonas (BR-AM)	Americas (Oregon)

Approximate write latency ⓘ
Low High



- 0% +

Query routing without LATAM zone

```
{  
  location:  
    {$in: [ 'BR-PE', 'BR-PA' ]}  
}
```

MongoDB Compass - globalcluster-okrrf.mongodb.net:27017/CustomerSingleView.Customers

Cluster SHARDED CLUSTER (9 MONGOSSES)

CustomerSingleView.Customers

DOCUMENTS 8 TOTAL 8 7.4

Documents Aggregations Schema Explain Plan Indexes

FILTER {location: {\$in: ['BR-PE', 'BR-PA']}}

VIEW DETAILS AS VISUAL TREE RAW JSON

Query Performance Summary

- Documents Returned: 2
- Index Keys Examined: 4
- Documents Examined: 2
- Actual Query Execution Time (ms): 165
- Sorted in Memory: no
- Query used the following index: location customerId

SHARD MERGE nReturned: 2 Execution Time: 165 ms

GLOBALCLUSTER-SHARD-2

GLOBALCLUSTER-SHARD-1

FETCH nReturned: 0 Execution Time: 0 ms

FETCH nReturned: 2 Execution Time: 0 ms

SHARDING FILTER nReturned: 0 Execution Time: 0 ms

SHARDING FILTER nReturned: 2 Execution Time: 0 ms

IXSCAN nReturned: 0 Execution Time: 0 ms

Index Name: location_1_customerId_1

Multi Key Index: no

IXSCAN nReturned: 2 Execution Time: 0 ms

Index Name: location_1_customerId_1

Multi Key Index: no

The screenshot shows the MongoDB Compass interface displaying the execution plan for a query. The plan starts with a 'SHARD MERGE' stage that returned 2 documents, taking 165ms. This stage is shown with a green circle highlighting the execution time. It branches into two parallel paths, one for 'GLOBALCLUSTER-SHARD-2' and one for 'GLOBALCLUSTER-SHARD-1'. Each path begins with a 'FETCH' stage (0ms) followed by a 'SHARDING FILTER' stage (0ms). These filter stages then lead to an 'IXSCAN' stage on the 'location_1_customerId_1' index (0ms). The total execution time for the query is 165ms.

Get ISO Country code from HTTP Request

AWS – Lambda & CloudFront

- `headers['cloudfront-viewer-country']`

*Limited to country

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/lambda-examples.html#lambda-examples-redirecting-examples>

Google Cloud – Functions

- `country: req.headers["x-appengine-country"]`
- `region: req.headers["x-appengine-region"]`

<https://cloud.google.com/appengine/docs/standard/go/reference/request-response-headers>

<https://medium.com/mop-developers/free-ip-based-geolocation-with-google-cloud-functions-f92e20d47651>

CloudFlare – Enable IP Geolocation

- `HTTP_CF_IPCOUNTRY` header

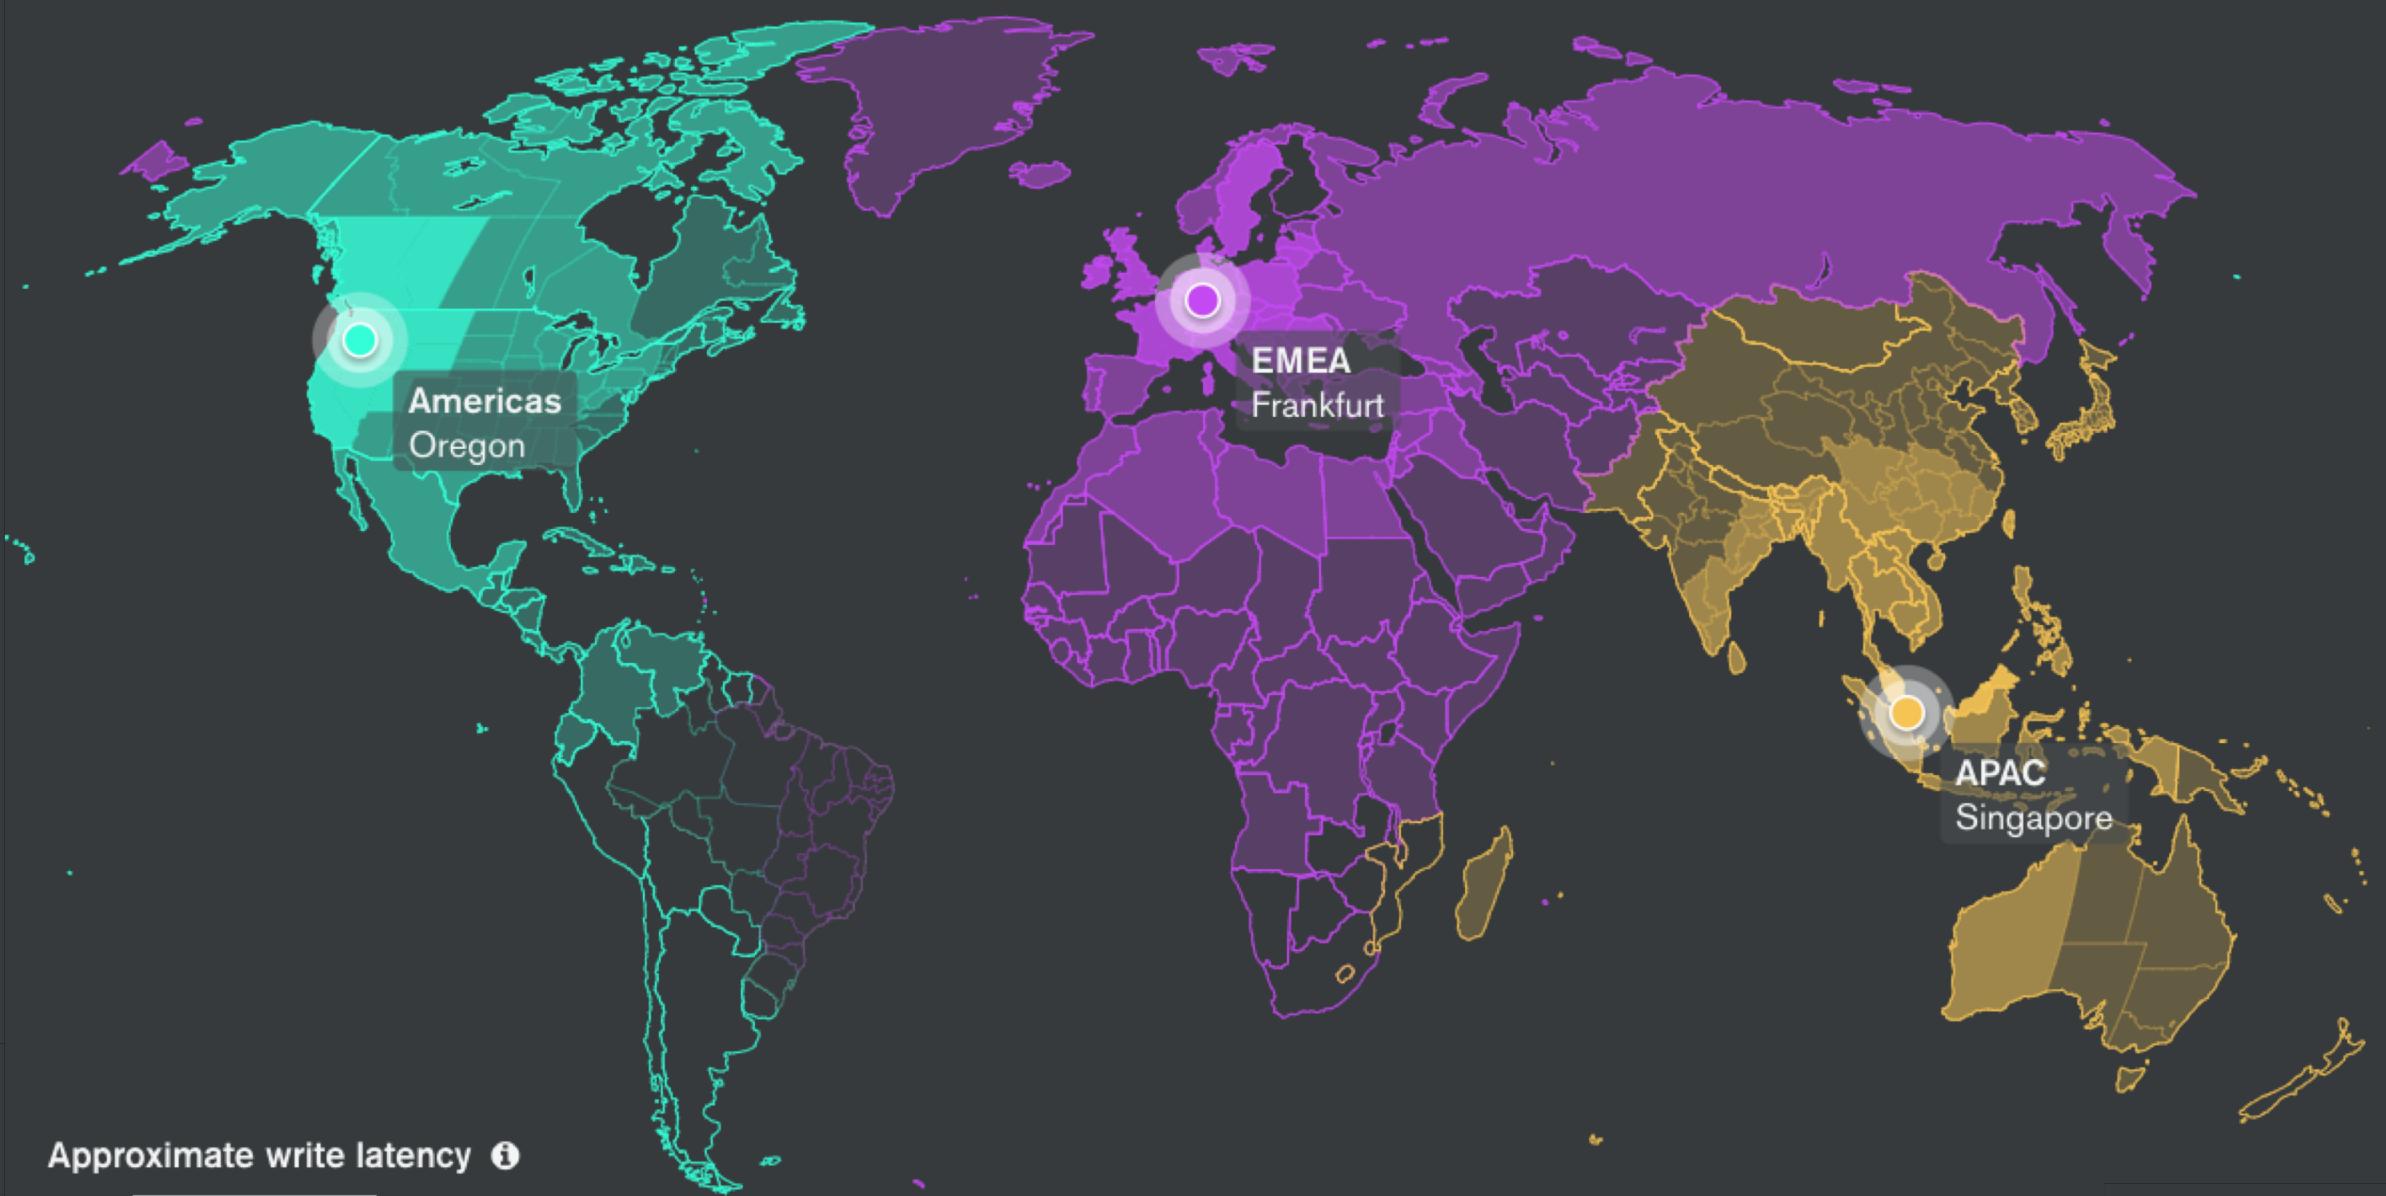
Akamai – EdgeScape (CDN)

- `country_code` and `region_code` headers



Provide a better experience to
growing markets

#MDBlocal



When should I add a Zone?

If latency of complete request exceeds threshold?

Not really – There are many layers contributing to the latency outside of Atlas

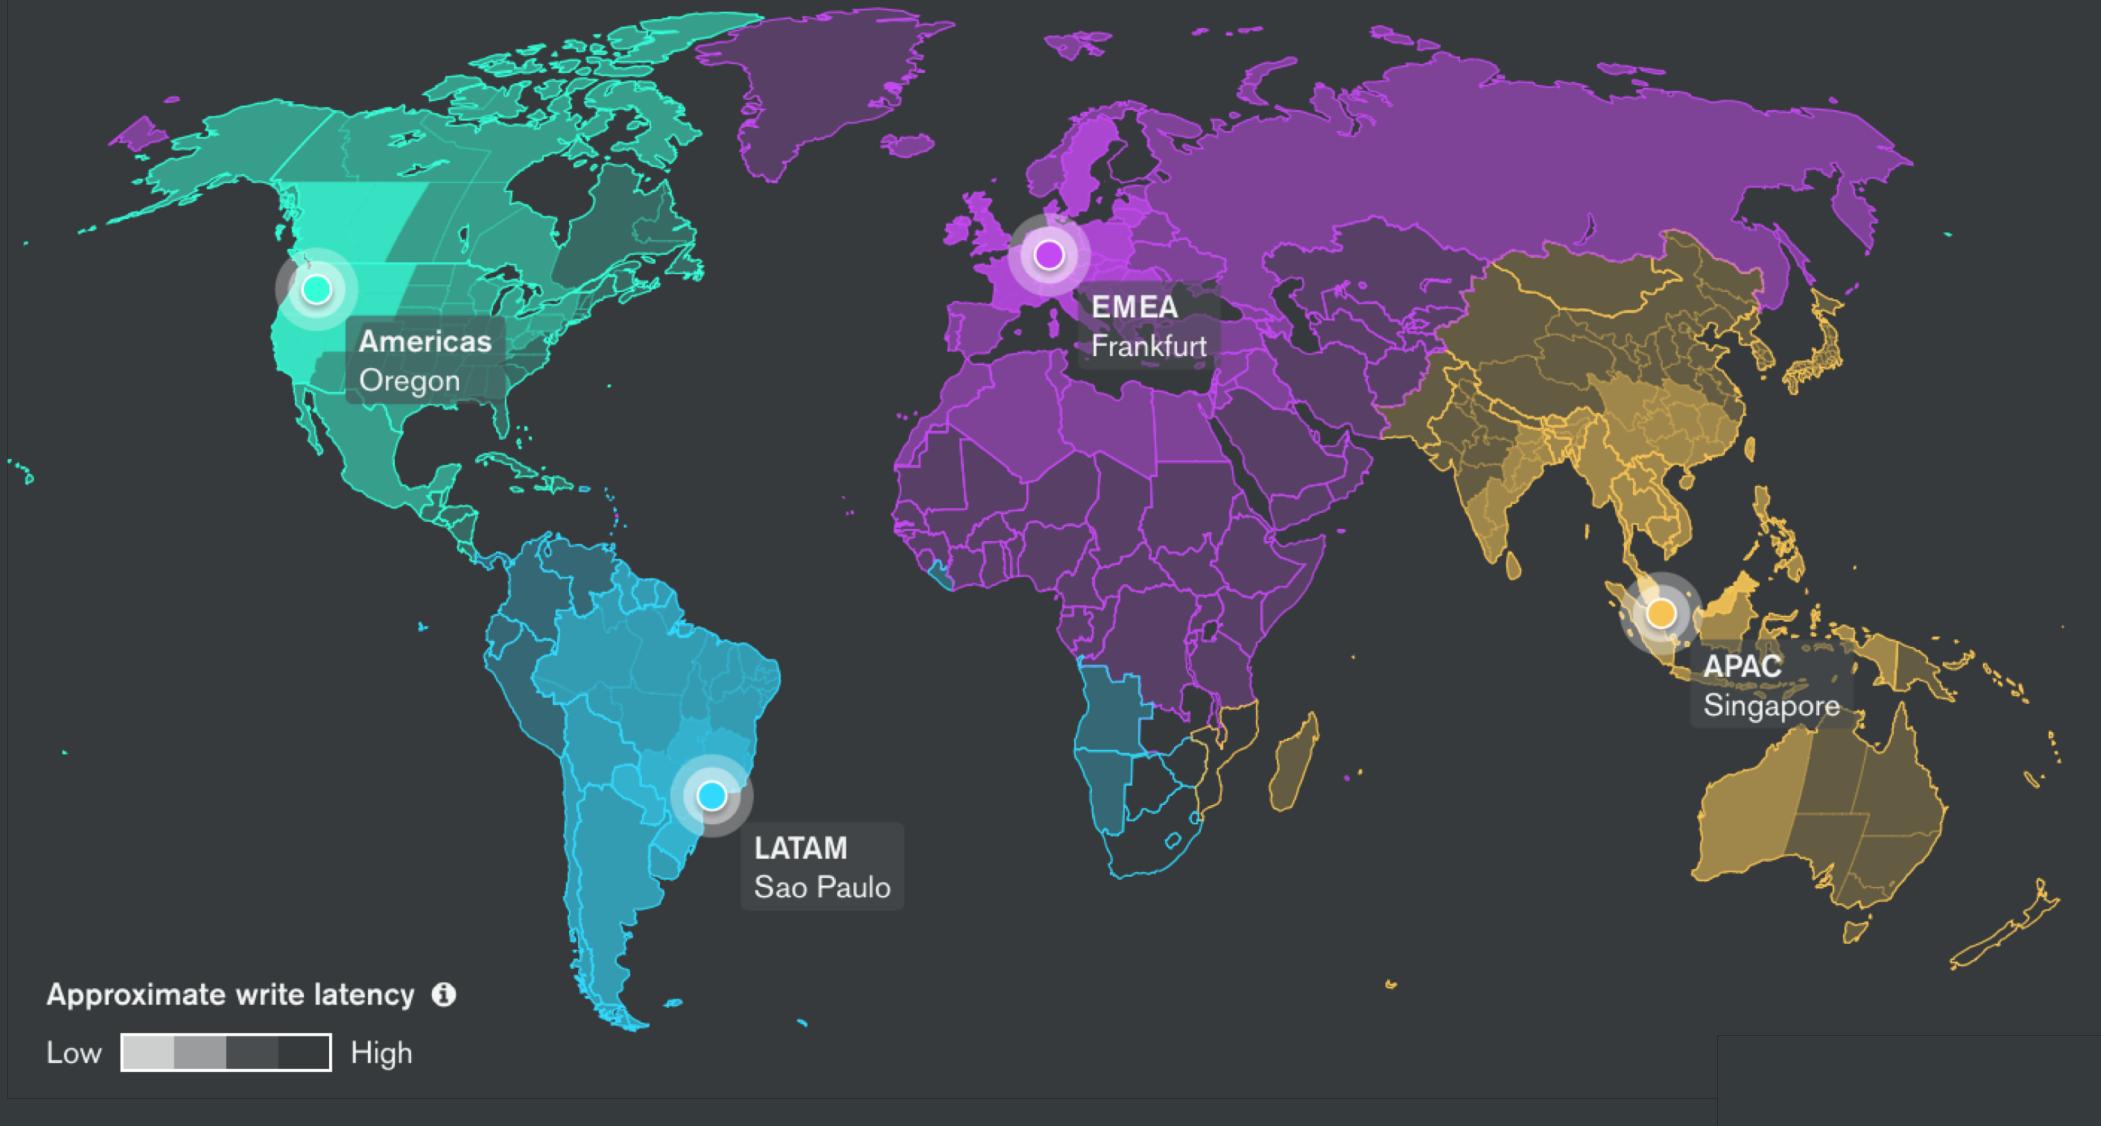
Metrics in Atlas?

Maybe – performance metrics

Maybe – mongos logs – cross-reference IP against IP-GeoLocation DB / Service

It's beyond Atlas - A combination of:

- Data Set growing from certain countries & sub-divisions
Use Aggregation Framework, BI Connector or Spark Connector for Analytics
- CDN / API Management Metrics: AWS CloudFront / API Gateway, Akamai, etc.
- Planned Expansion: New features targeting specific markets



Query routing with LATAM zone

```
{  
  location:  
    {$in: [ 'BR-PE', 'BR-PA' ]}  
}
```

MongoDB Compass - globalcluster-breakoutsession-twigg.mongodb.net:27017/CustomerSingleView

My Cluster

4 DBS 14 COLLECTIONS

filter

CustomerSingleView

Customers

admin config local

Cluster SHARDED CLUSTER 9 MONGOSSES

CustomerSingleView.Customers

Documents Aggregations Schema Explain Plan Indexes

FILTER {location: {\$in: ['BR-PE', 'BR-PA']}}

VIEW DETAILS AS VISUAL TREE RAW JSON

Query Performance Summary

- Documents Returned: 1868
- Index Keys Examined: 1870
- Documents Examined: 1868

Actual Query Execution Time (ms): 169

Sorted in Memory: no

Query used the following index:
location ↗ customerId ↗

SINGLE_SHARD
nReturned: 1868 Execution Time: 169 ms

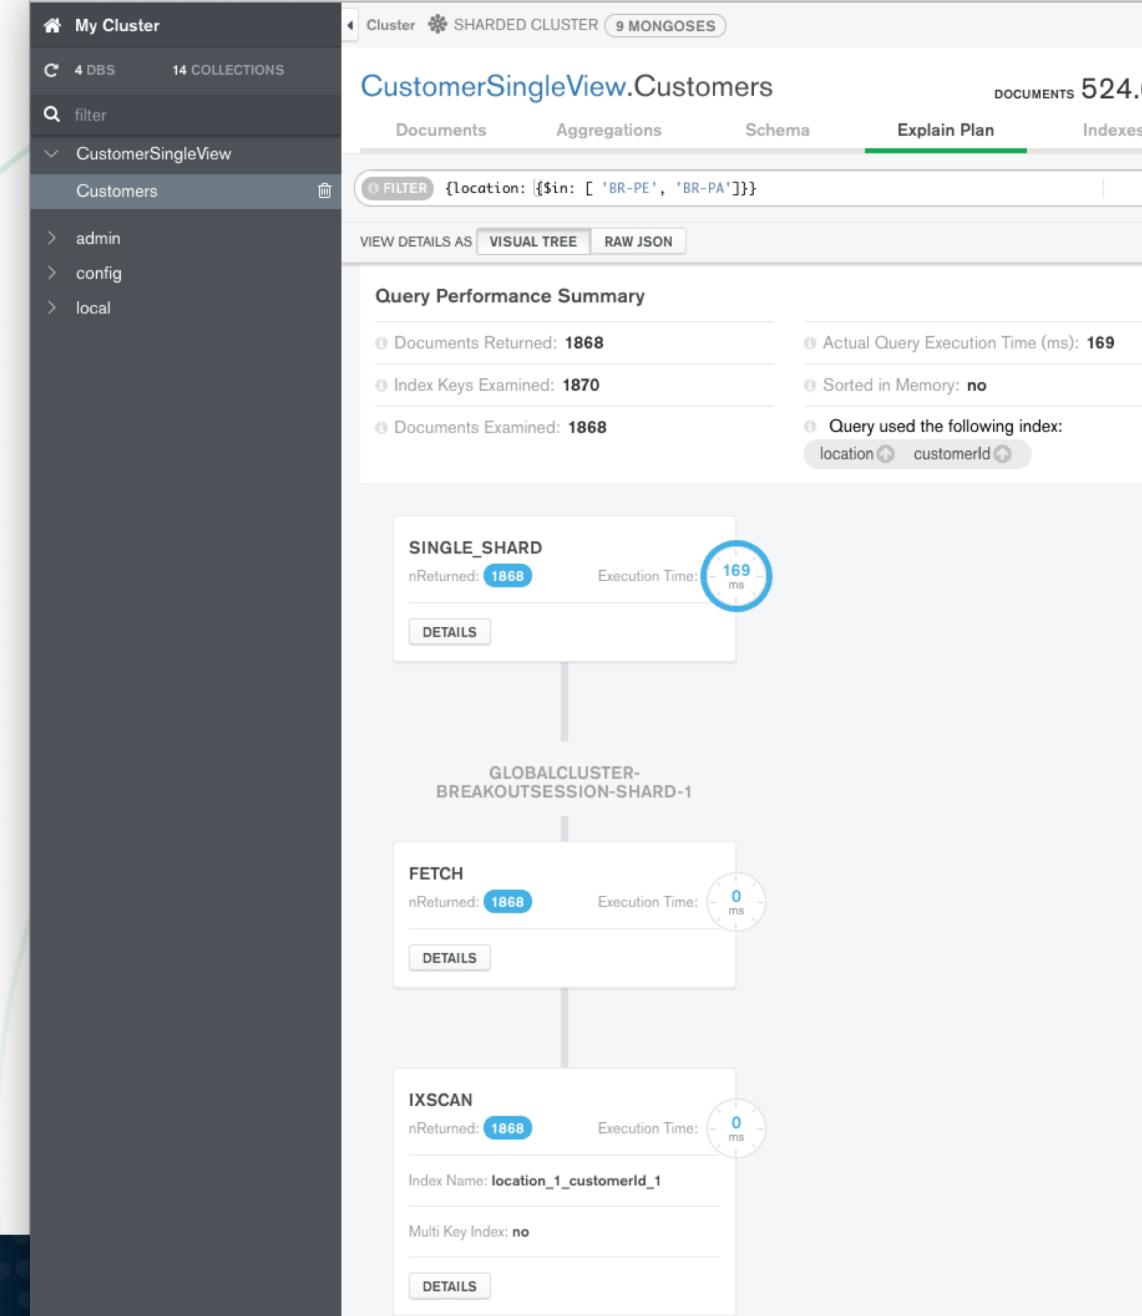
GLOBALCLUSTER-BREAKOUTSESSION-SHARD-1

FETCH
nReturned: 1868 Execution Time: 0 ms

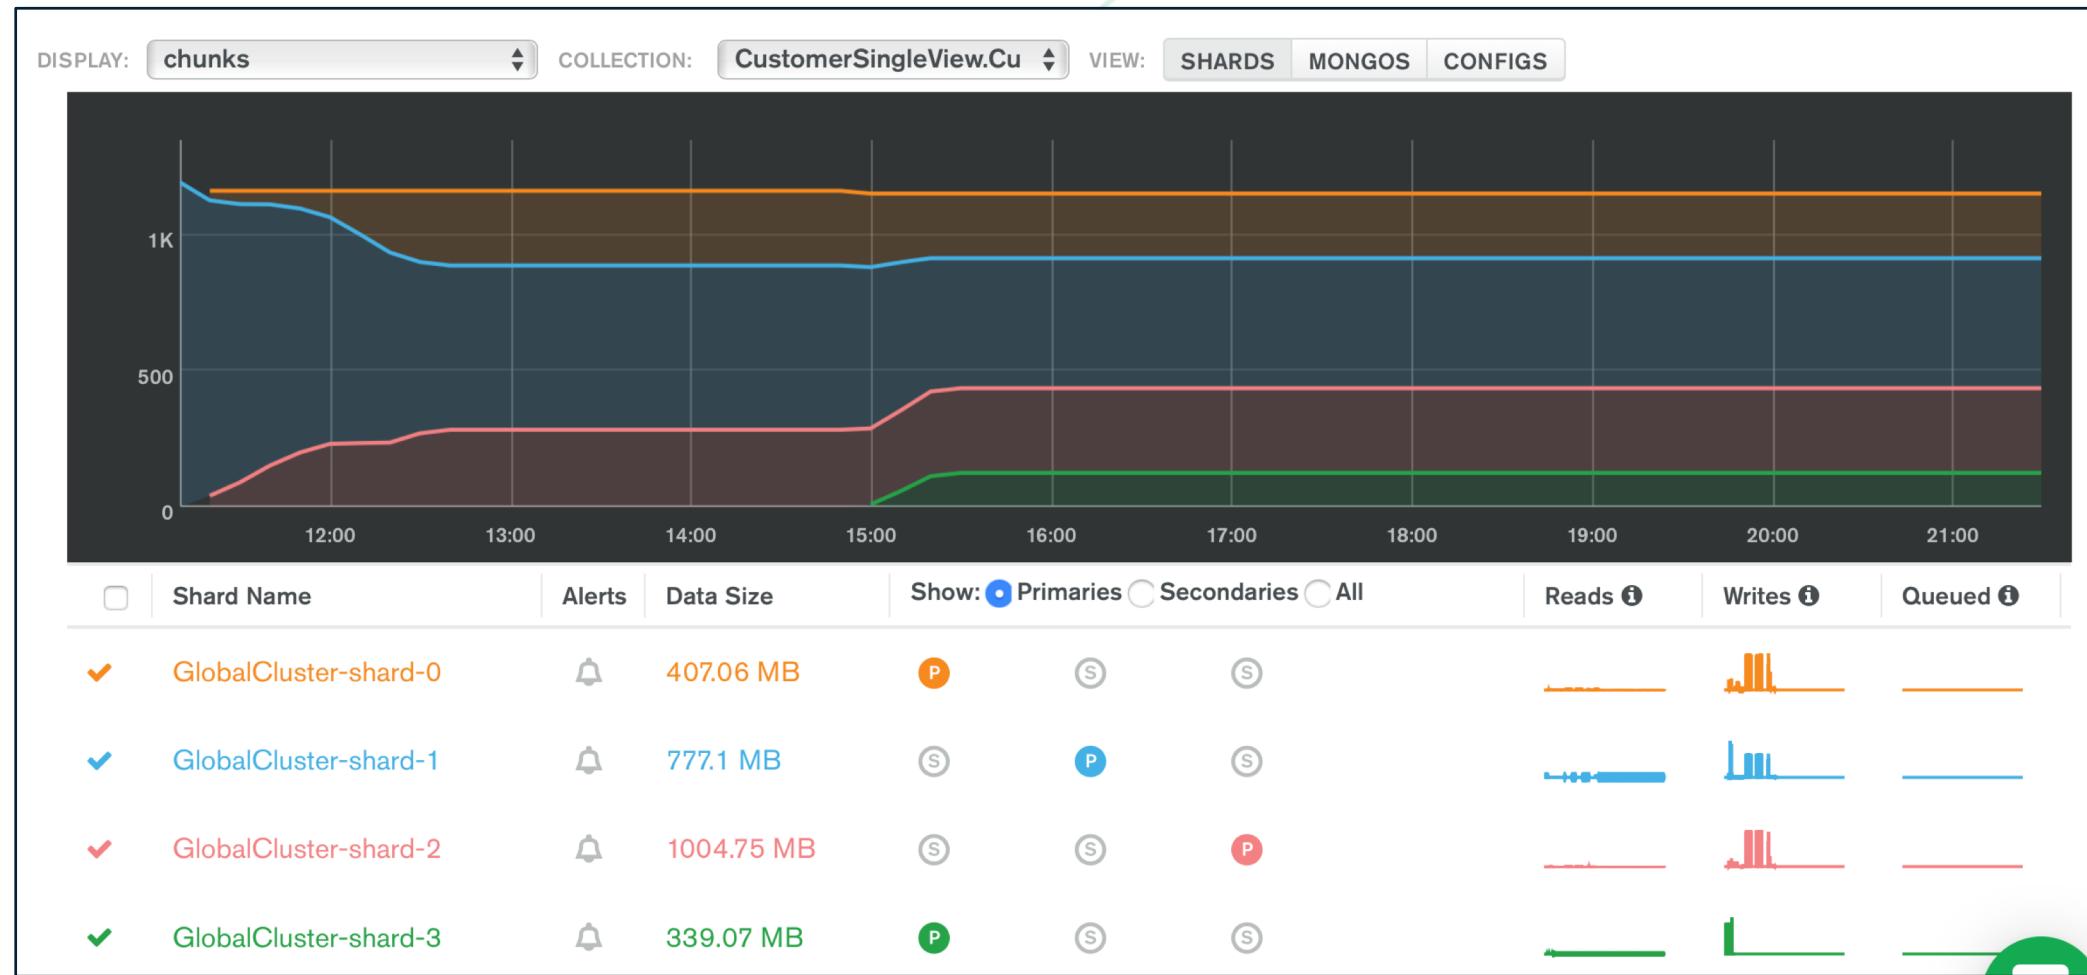
IXSCAN
nReturned: 1868 Execution Time: 0 ms

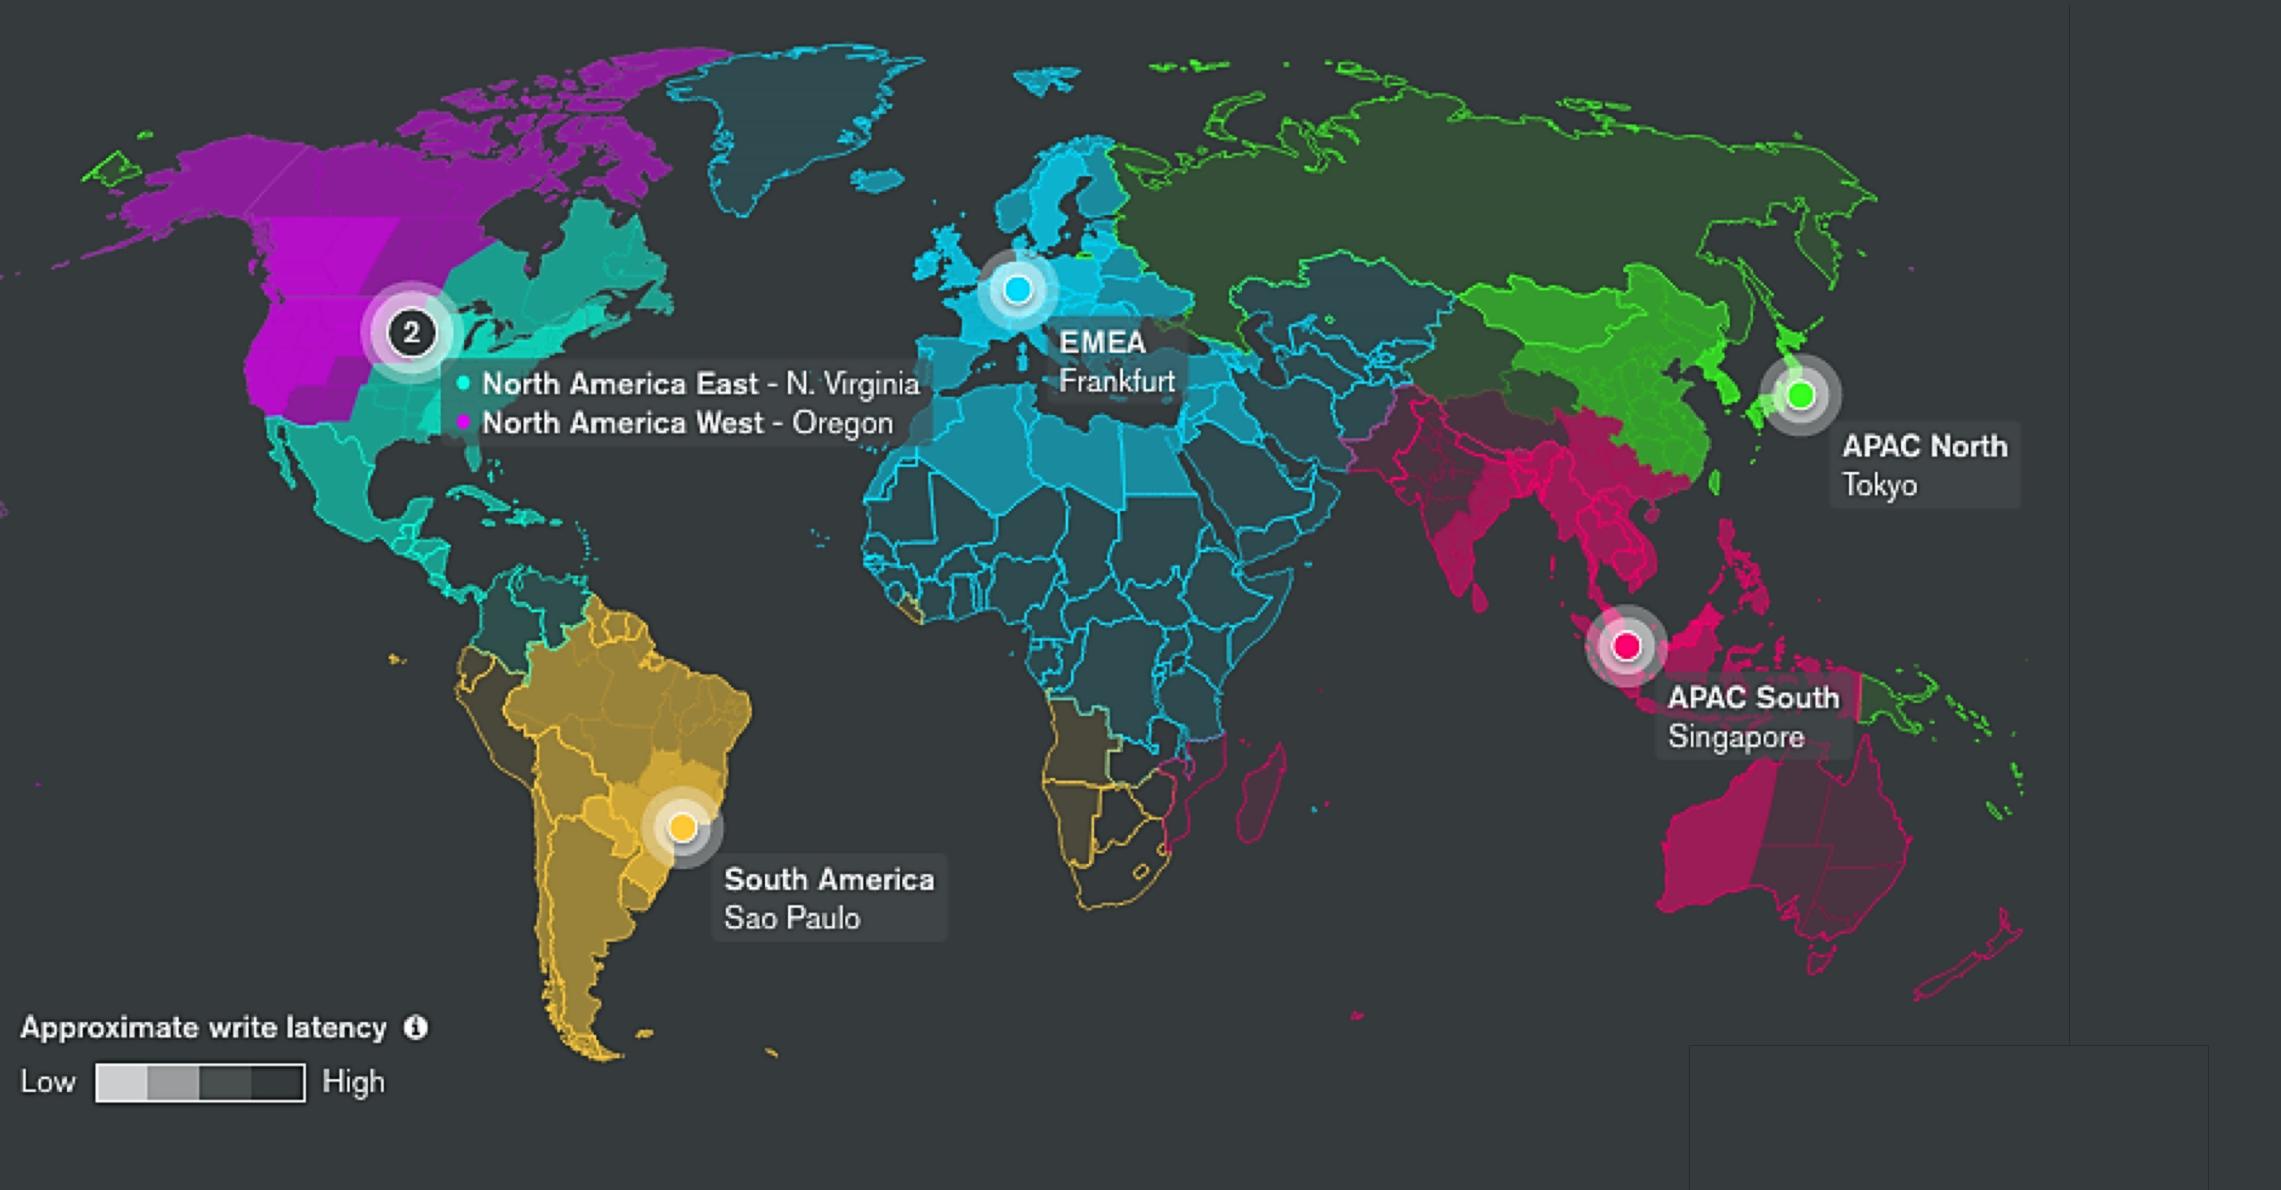
Index Name: location_1_customerId_1

Multi Key Index: no



Chunk Migration – adding LATAM zone





NA West

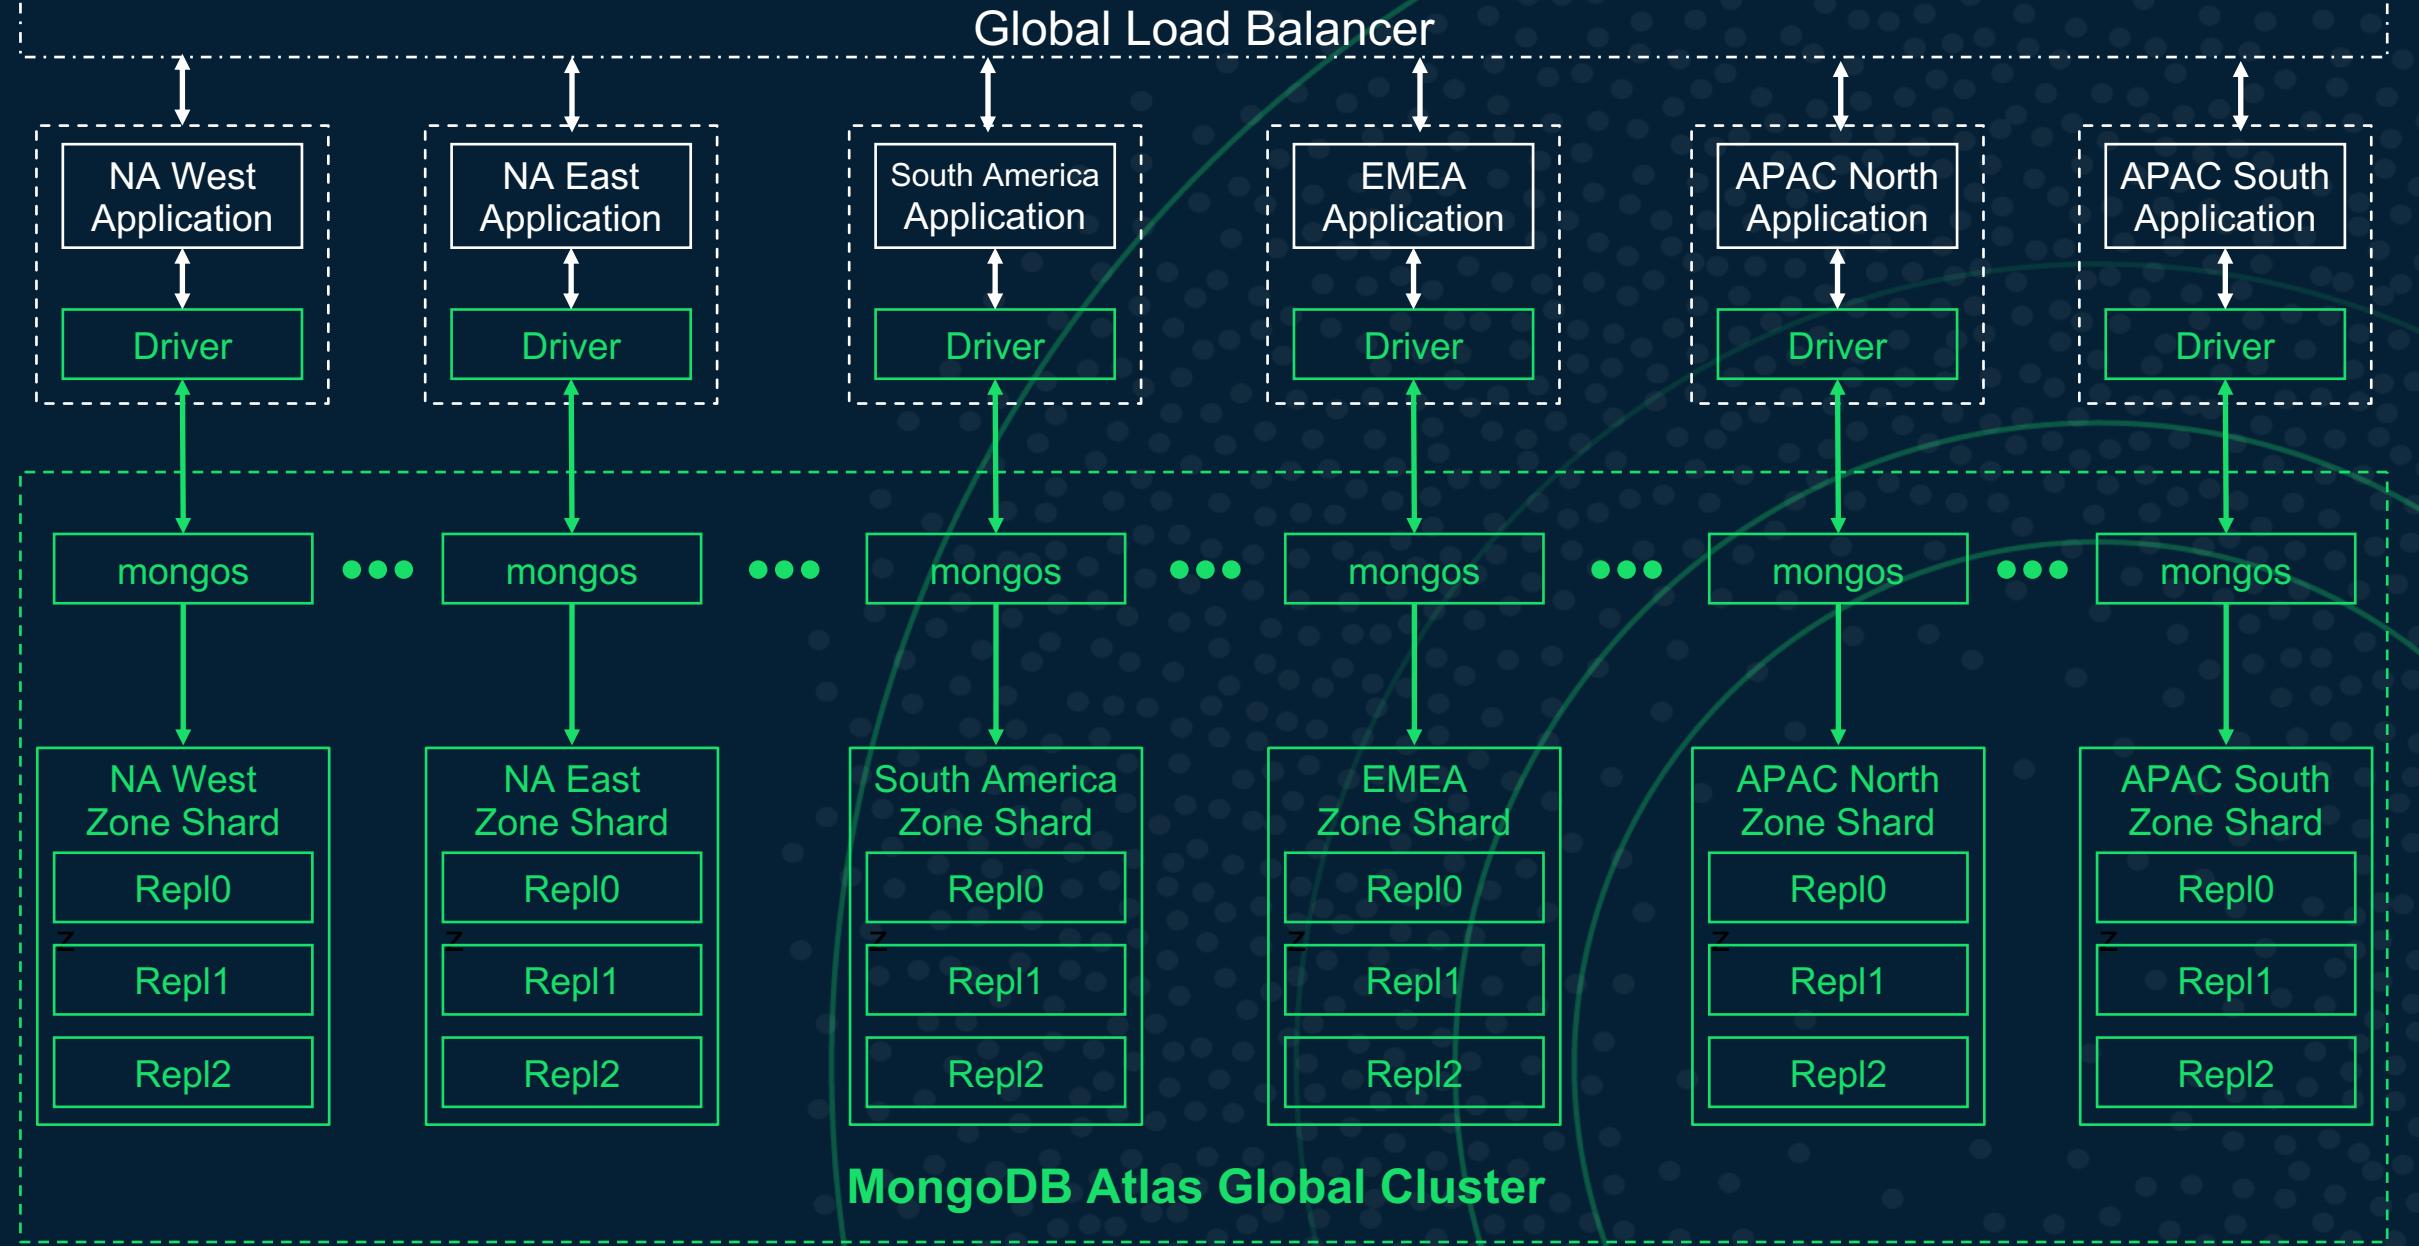
NA East

South America

EMEA

APAC North

APAC South



Simplified connectivity

#MDBlocal

Concise and global, it will look like

```
mongodb+srv://user:pass@globalCluster-rzeyq.mongodb.net
```

Your driver will use this address to discover every mongos in the cluster globally

Taking advantage of the MongoDB 3.6 driver support for SRV records

... so you can deploy the same application code everywhere - symmetrically

Resolving SRV record to mongos's

Blog Post: <https://www.mongodb.com/blog/post/mongodb-3-6-here-to-SRV-you-with-easier-replica-set-connections>

```
$ python mongodb_srv_records.py globalcluster-twijg.mongodb.net  
globalcluster-shard-00-00-twijg.mongodb.net:27016  
globalcluster-shard-00-01-twijg.mongodb.net:27016  
globalcluster-shard-00-02-twijg.mongodb.net:27016  
globalcluster-shard-01-00-twijg.mongodb.net:27016  
globalcluster-shard-01-01-twijg.mongodb.net:27016  
globalcluster-shard-01-02-twijg.mongodb.net:27016  
globalcluster-shard-02-00-twijg.mongodb.net:27016  
globalcluster-shard-02-01-twijg.mongodb.net:27016  
globalcluster-shard-02-02-twijg.mongodb.net:27016  
globalcluster-shard-03-00-twijg.mongodb.net:27016  
globalcluster-shard-03-01-twijg.mongodb.net:27016  
globalcluster-shard-03-02-twijg.mongodb.net:27016  
globalcluster-twijg.mongodb.net: "authSource=admin"
```

An even better user experience?

What if I want low latency global reads?

What if I want low latency global reads?

Zone configuration summary

→ CONFIGURE LOCAL READS IN ALL ZONES ⓘ

● Americas	N. Virginia (us-east-1)	<ul style="list-style-type: none">✓ Low latency reads and writes in North America✗ Local reads in all other zones✓ Available during partial region outage✗ Not available during full region outage
● EMEA	Frankfurt (eu-central-1)	<ul style="list-style-type: none">✓ Low latency reads and writes in Europe✗ Local reads in all other zones✓ Available during partial region outage✗ Not available during full region outage
● APAC	Singapore (ap-southeast-1)	<ul style="list-style-type: none">✓ Low latency reads and writes in Asia✗ Local reads in all other zones✓ Available during partial region outage✗ Not available during full region outage

What if I want low latency global reads?

- ✓ Deploy across multiple regions ⓘ

Node Type	Region	Number of Nodes
Preferred ⓘ	 Oregon (us-west-2)	3
Add a region		

- ✓ Deploy read-only replicas ⓘ

Node Type	Region	Number of Nodes	
Read-only	 Frankfurt (eu-central-1)	1	
Read-only	 Singapore (ap-southeast-1)	1	
Read-only	 Sao Paulo (sa-east-1)	1	

Where this is going and limits

Massive reach

MongoDB Atlas

- in-country regions for 27% of the world population
- with a much higher percent served by regions in nearby countries

Available in 56 global cloud regions

- AWS: 15 regions (38 availability zones)
- Azure: 26 regions (26 availability zones w/ 60 Fault Domains)
- GCP: 15 regions (38 Availability zones)

There are gaps

Africa

- S. Africa announced by Azure; a long way from most of Africa's population, however
- Biggest economies in Africa: Nigeria, South Africa, Egypt; with Ethiopia now the fastest growing

Middle East

- We'll see

China

- It's complicated...

Russia

- It's also complicated...
- Frankfurt remains best bet but also Finland on Google Cloud

Special / **Government** / Compliance contexts

But MongoDB runs everywhere

Over time the entire world will be well served by major cloud infrastructure

You can always run your own MongoDB outside of where the big public cloud

- MongoDB Ops Manager / MongoDB Cloud Manager
- Live Migrate to MongoDB Atlas when ready

So you can have a seamless development experience in any context

Imagine 2030

2030 Projection

GDP & relative position
compared with today



You can go global incrementally

Path to Global Over Time

1. **Cluster – Single Region** ← deploy today for free on Atlas
 - Focus on finding traction
 - Scale up when needed in a few minutes on Atlas

Path to Global Over Time

1. **Cluster – Single Region** ← deploy today for free on Atlas
 - Focus on finding traction
 - Scale up when needed in a few minutes on Atlas
2. **Cluster – Multi Region HA**
 - Get your application HA as well
 - Scale up when needed in a few minutes on Atlas

Path to Global Over Time

1. **Cluster – Single Region** ← deploy today for free on Atlas
 - Focus on finding traction
 - Scale up when needed in a few minutes on Atlas
2. **Cluster – Multi Region HA**
 - Get your application HA as well
 - Scale up when needed in a few minutes on Atlas
3. **Cluster – add Global Read Regions**
 - Route users to nearby application instance; use secondary reads
 - Scale up when needed in a few minutes on Atlas

Path to Global Over Time

1. **Cluster – Single Region** ← deploy today for free on Atlas
 - Focus on finding traction
 - Scale up when needed in a few minutes on Atlas
2. **Cluster – Multi Region HA**
 - Get your application HA as well
 - Scale up when needed in a few minutes on Atlas
3. **Cluster – add Global Read Regions**
 - Route users to nearby application instance; use secondary reads
 - Scale up when needed in a few minutes on Atlas
4. **Enable Global Writes**
 - Add a second write zone
 - Add location attribute to your data & shard your global collections in the Atlas UI

Path to Global Over Time

1. **Cluster – Single Region** ← deploy today for free on Atlas
 - Focus on finding traction
 - Scale up when needed in a few minutes on Atlas
2. **Cluster – Multi Region HA**
 - Get your application HA as well
 - Scale up when needed in a few minutes on Atlas
3. **Cluster – add Global Read Regions**
 - Route users to nearby application instance; use secondary reads
 - Scale up when needed in a few minutes on Atlas
4. **Enable Global Writes**
 - Add a second write zone
 - Add location attribute to your data & shard your global collections in the Atlas UI
5. **Expand Reach of Global Cluster**
 - Add more write zones
 - Or add shards to existing zones
 - Scale up any time

And you can follow your users...

**Now go out and reach
those 4+ billion internet users**



Thanks

Sigfrido “Sig” Narvaez

Principal Solution Architect, MongoDB

sig@mongodb.com





Our developer focused talks are back on the road!

Find one near you

At your MongoDB.local, you'll learn **technologies**, **tools**, and **best practices** that make it easy for you to build data-driven applications without distraction.

mongodb.com/local

#MDBlocal